

## Table of contents

|                                                |          |
|------------------------------------------------|----------|
| <b>The Essentials</b>                          | <b>1</b> |
| Why Study High-Dimensional Spaces? . . . . .   | 1        |
| The Representation Space . . . . .             | 2        |
| The Curse of Dimensionality . . . . .          | 2        |
| 1. Data Sparsity . . . . .                     | 2        |
| 2. Empty Sphere and Spiky Cube . . . . .       | 3        |
| 3. Volume Distribution . . . . .               | 4        |
| 4. The Gaussian Egg . . . . .                  | 5        |
| 5. Distance Concentration . . . . .            | 5        |
| 6. Quasi-Orthogonality . . . . .               | 6        |
| 7. Hubness . . . . .                           | 7        |
| 8. Concentration Inequalities . . . . .        | 8        |
| The Blessing of Dimensionality . . . . .       | 9        |
| 1. Laws of Large Numbers . . . . .             | 9        |
| 2. Random Projections . . . . .                | 10       |
| Impact in Data Science . . . . .               | 11       |
| 1. Sampling . . . . .                          | 11       |
| 2. Indexing . . . . .                          | 11       |
| 3. Learning . . . . .                          | 12       |
| 4. Why Does Data Science Still Work? . . . . . | 12       |
| Global Summary . . . . .                       | 13       |
| Positive Aspects (Blessing) . . . . .          | 13       |
| Negative Aspects (Curse) . . . . .             | 13       |
| Unique Solution . . . . .                      | 14       |

## The Essentials

### Why Study High-Dimensional Spaces?

**Data Science** studies how to represent the real world in the form of numerical data within mathematical spaces. It focuses on the **geometric** and **statistical** properties of these spaces.

#### Universal Pipeline:

- Real world  $\rightarrow$  Sensors  $\rightarrow$  Raw data  $\rightarrow$  Numerical representation
- Everything becomes a **matrix**:  $N$  samples,  $D$  features

#### Two Complementary Perspectives:

- **Statistical**: Sample from an underlying distribution

- **Algebraic:** Matrix  $X \in \mathbb{R}^{D \times N}$
- 

## The Representation Space

A  $D$ -dimensional space has a complete structure:

- Inner product  $\rightarrow$  Norm  $\rightarrow$  Distance  $\rightarrow$  Metric space

**Central Question:** What happens when  $D$  becomes large?

- Negative effects: **Curse** of dimensionality
  - Positive effects: **Blessing** of dimensionality
- 

## The Curse of Dimensionality

### 1. Data Sparsity

To accurately estimate a distribution in  $D$  dimensions, an **exponential** number of samples is required:

$$N \approx n \cdot b^D$$

where  $n$  = samples per bin,  $b$  = number of bins per dimension.

**Example:**

- $D = 6$ ,  $b = 10$ ,  $n = 10 \rightarrow N \approx 10^7$  samples needed
- With fixed  $N$ : most bins remain **empty**

#### Details: Estimation Quality

If  $N$  is fixed, the estimation quality becomes:

$$n = \frac{N}{b^D} \implies \mathbb{E}[P(x_i \in \text{bin}_j)] \approx \frac{1}{b^D}$$

The probability that a bin is occupied decreases exponentially with  $D$ .

### Quick Review: Data Sparsity

Requires  $N \approx n \cdot b^D$  samples (**exponential** growth).

If  $N$  is fixed:  $n = N/b^D \rightarrow$  most bins **empty**.

## 2. Empty Sphere and Spiky Cube

In a high-dimensional cube, the volume of the inscribed sphere becomes **negligible**:

$$\frac{\text{vol}(B_D(r))}{\text{vol}(C_D(r))} = \frac{\pi^{D/2}}{D \cdot 2^{D-1} \Gamma(D/2)} \xrightarrow{D \rightarrow \infty} 0$$

**Consequences:**

- Volume concentrates in the **corners** of the cube
- Center-corner distance  $= \frac{\sqrt{D}}{2}$  (grows with  $\sqrt{D}$ )
- Uniform sampling  $\rightarrow$  **empty** sphere

### Details: Volume of the Hypersphere

**General Formula:**

$$\text{vol}(B_D(r)) = \frac{2r^D \pi^{D/2}}{D \Gamma(D/2)}$$

**Recurrence Relation** (for  $D > 1$ ):

$$\text{vol}(B_D(r)) = \frac{2\pi r^2}{D} \text{vol}(B_{D-2}(r))$$

**Corners of the Unit Cube:**

- Number of corners:  $2^D$
- Distance to origin:  $\|x_i\|_2 = \frac{\sqrt{D}}{2}$

| $D$                  | 1             | 2                    | 4 | 5                    | 100 |
|----------------------|---------------|----------------------|---|----------------------|-----|
| $\frac{\sqrt{D}}{2}$ | $\frac{1}{2}$ | $\frac{\sqrt{2}}{2}$ | 1 | $\frac{\sqrt{5}}{2}$ | 5   |

### Quick Review: Empty Sphere & Spiky Cube

Sphere/cube ratio  $\rightarrow 0$ . Volume concentrated at **corners** (distance  $\sqrt{D}/2$ ).  
Uniform sampling  $\rightarrow$  **empty** sphere.

---

### 3. Volume Distribution

For **any shape** (cube, sphere, or other), volume concentrates near the surface:

$$\lim_{D \rightarrow \infty} \frac{\text{vol}(V_D(r)) - \text{vol}(V_D((1-\varepsilon)r))}{\text{vol}(V_D(r))} = 1$$

**Principle:**

- Reducing size by factor  $(1-\varepsilon) \rightarrow$  volume reduced by  $(1-\varepsilon)^D \rightarrow 0$
- Almost all volume is in a thin **shell** at the surface

#### Details: Calculation for Hypercube and Hypersphere

**For the Hypercube:**

$$\frac{\text{vol}(C_D((1-\varepsilon)r))}{\text{vol}(C_D(r))} = (1-\varepsilon)^D \xrightarrow{D \rightarrow \infty} 0$$

**For the Hypersphere:**

$$\frac{\text{vol}(B_D((1-\varepsilon)r))}{\text{vol}(B_D(r))} = (1-\varepsilon)^D \xrightarrow{D \rightarrow \infty} 0$$

**Technical Note:**  $(1-\varepsilon)^D \leq e^{-D\varepsilon} \rightarrow 0$  as  $D \rightarrow \infty$

**Decomposition Method:** Divide the volume into arbitrarily small hypercubic voxels, then apply the same reasoning.

### Quick Review: Volume Distribution

For **any shape:** ratio  $(1-\varepsilon)^D \rightarrow 0$ .  
Volume **concentrated at the surface**.

---

#### 4. The Gaussian Egg

For a distribution  $X \sim \mathcal{N}(0, \text{Id}_D)$ , the radius  $R = \|X\|_2$  follows:

$$R^2 = \sum_{i=1}^D X_i^2 \sim \chi^2(D) \implies \mathbb{E}[R^2] = D \implies \mathbb{E}[R] \approx \sqrt{D}$$

The density is **concentrated** around the radius  $\sqrt{D}$ .

**Gaussian Annulus Theorem:**

$$P\left[\left|\|X\|_2 - \sqrt{D}\right| > \beta\right] \leq 3e^{-c\beta^2}$$

Almost all density is in a **thin ring** around  $\sqrt{D}$ .

##### Details: Density Along the Radius

**Context:** Coordinates  $X_i \sim \mathcal{N}(0, 1)$  i.i.d.

**Integration along  $r$**  = take the norm.  $R$  is the random variable of the radius.

**Alternative Formulation:**

$$P\left[\left|\frac{\|X\|}{\sqrt{D}} - 1\right| \leq \varepsilon\right] \approx 1$$

For a Gaussian  $\mathcal{N}(0, \sigma^2 \text{Id}_D)$ , the density is concentrated at radius  $\sigma\sqrt{D}$ .

##### Quick Review: Gaussian Egg

For  $X \sim \mathcal{N}(0, \text{Id}_D)$ :  $R^2 \sim \chi^2(D)$  so  $\mathbb{E}[R^2] = D$ .

Density concentrated at radius  $\sqrt{D}$  (**thin ring**).

---

#### 5. Distance Concentration

All neighbors become **equidistant** as  $D$  increases.

**Theorem:** If the variance of the normalized distance tends to 0, then:

$$\lim_{D \rightarrow \infty} P\left[\frac{D_{\max} - D_{\min}}{D_{\min}} \leq \varepsilon\right] = 1$$

where  $D_{\min}$  and  $D_{\max}$  are the distances to the nearest and farthest neighbors.

### Consequences:

- All neighbors appear at distance  $D_{\min}$
- Discriminative power decreases exponentially
- KNN and  $\varepsilon$ -NN become **non-significant**

#### Details: Concentration Condition

**Complete Theorem:** If  $\lim_{D \rightarrow \infty} \text{Var} \left[ \frac{d(x,y)^p}{\mathbb{E}[d(x,y)^p]} \right] = 0$  for  $0 < p < \infty$ , then concentration occurs.

This condition means that the relative fluctuations of the distance become negligible.

#### Quick Review: Distance Concentration

If normalized distance variance  $\rightarrow 0$ , then  $\frac{D_{\max} - D_{\min}}{D_{\min}} \rightarrow 0$ .  
All neighbors **equidistant**. KNN loses meaning.

---

## 6. Quasi-Orthogonality

In high dimensions, random vectors exhibit surprising geometric properties:

**Two random vectors** are quasi-orthogonal:

$$\mathbb{E}[\cos(\theta)] = 0 \implies \theta \approx \frac{\pi}{2}$$

**Random triangle** is quasi-equilateral:

$$\mathbb{E}[\cos(\theta)] = \frac{1}{2} \implies \theta \approx \frac{\pi}{3}$$

**Consequence:** On a sphere, volume concentrates at the **equator** (regardless of direction).

#### Details: Angle Calculation Between Vectors

**For two random vectors:**

Define 4 i.i.d. random variables  $W, X, Y, Z$  on  $\Omega \subset \mathbb{R}^D$ :

$$\mathbb{E}[\cos(\angle(X - Y, Z - W))] = \mathbb{E} \left[ \frac{\langle X - Y, Z - W \rangle}{\|X - Y\| \cdot \|Z - W\|} \right]$$

Since  $W, X, Y, Z$  are i.i.d.,  $X - Y$  and  $Z - W$  are also i.i.d. and centered:  $\mathbb{E}(X - Y) = \mathbb{E}(Z - W) = 0$ .

The numerator is  $\text{cov}(X - Y, Z - W) = 0$ .

**For a triangle:**

With 3 i.i.d. random variables  $X, Y, Z$ , compute:

$$\frac{\langle X - Y, Z - Y \rangle}{\langle X - Y, X - Y \rangle} = \frac{\langle X, Z \rangle - \langle X, Y \rangle - \langle Y, Z \rangle + \langle Y, Y \rangle}{\langle X, X \rangle + \langle Y, Y \rangle - 2\langle X, Y \rangle}$$

In expectation:  $= \frac{1}{2}$

**Note:** To sample on the unit hypersphere, sample a centrally symmetric distribution (e.g., Gaussian) and normalize.

### Quick Review: Quasi-Orthogonality

**2 random vectors:**  $\theta \approx \pi/2$  (quasi-orthogonal).

**Random triangle:**  $\theta \approx \pi/3$  (equilateral).

**Sphere:** volume at the equator.

---

## 7. Hubness

Some samples appear in the neighborhood of **many** others.

**Definition:**  $n_k(x_j)$  = number of times  $x_j$  is a k-NN of another point.

**Observation:** The distribution of  $n_k$  is **asymmetric**:

- A few points (**hubs**) have very high values
- Most points have low values
- These hubs are “close” to all other samples

### Details: Formal Definition of Hubness

Given  $X$ , define:

$$r_{ik}(x_j) = \begin{cases} 1 & \text{if } x_j \in V_X^k(x_i) \\ 0 & \text{otherwise} \end{cases}$$

Then:

$$n_k(x_j) = \sum_i r_{ik}(x_j)$$

**Empirical Visualization:** For 20-NN with  $D = 100$  and 1000 uniform samples on 50 bins, a strong asymmetry in the distribution of  $n_k$  is observed.

### Quick Review: Hubness

$n_k(x_j)$  = number of times  $x_j$  is a k-NN of another point.

**Asymmetric** distribution: few **hubs** in many neighborhoods.

## 8. Concentration Inequalities

Concentration inequalities mathematically explain these phenomena.

### Fundamental Inequalities:

- **Markov:**  $P(X > \varepsilon) \leq \frac{\mathbb{E}[X]}{\varepsilon}$
- **Chebyshev:**  $P(|X - \mu| > \varepsilon) \leq \frac{\sigma^2}{\varepsilon^2}$

**Hoeffding's Inequality:** For  $\{X_i\}_{i \in [D]}$  independent with  $P(a \leq X_i \leq b) = 1$  and  $\mathbb{E}[X_i] = \mu$ :

$$P(|\bar{X} - \mu| > \varepsilon) \leq 2e^{-2D\varepsilon^2/(b-a)^2}$$

where  $\bar{X} = \frac{1}{D} \sum_{i=1}^D X_i$

### Explanation:

- Aggregation of independent variables  $\rightarrow$  convergence to  $\mu$
- Minkowski distances have structure  $\sum_{i=1}^D (\cdot)$
- For  $d_p(x, y)$  to be small (large), **all**  $D$  coordinates must be similar (dissimilar)
- As  $D$  increases, this becomes unlikely

### Details: Application to Distances

Minkowski distances:

$$d_p(x, y) = \left( \sum_{i=1}^D |x[i] - y[i]|^p \right)^{1/p}$$

Hoeffding's inequality formally explains why distance values **concentrate**.

**Intuition:** For large  $D$ , extreme values of  $d_p(x, y)$  become unlikely because they require coordination across **all** dimensions.

**Note:** Example of PAC (Probably Approximately Correct) analysis.



### Quick Review: Concentration Inequalities

**Hoeffding:**  $P(|\bar{X} - \mu| > \varepsilon) \leq 2e^{-2D\varepsilon^2/(b-a)^2}$ .  
Structure  $\sum_{i=1}^D (\cdot)$  in distances  $\rightarrow$  **concentration**.

## The Blessing of Dimensionality

### 1. Laws of Large Numbers

#### Weak Law of Large Numbers (WLLN):

For  $\{X_i\}_{i \in [D]}$  i.i.d. with  $\mathbb{E}[X_i] = \mu$ :

$$\lim_{D \rightarrow \infty} P(|\bar{X} - \mu| < \varepsilon) = 1$$

where  $\bar{X} = \frac{1}{D} \sum_{i=1}^D X_i$

#### Interpretation:

- The empirical mean  $\bar{X}$  converges to the true expectation  $\mu$
- Fundamental basis of estimation theory
- Consistent with the frequentist approach to probabilities

#### Central Limit Theorem (CLT):

$$Z_D = \frac{\sqrt{D}}{\sigma}(\bar{X} - \mu) \sim \mathcal{N}(0, 1)$$

if  $\text{Var}(X_i) = \sigma^2$

#### Details: Proof via Chebyshev

The WLLN can be proven from Chebyshev's inequality:

$$P(|\bar{X} - \mu| \geq \varepsilon) \leq \frac{\text{Var}(\bar{X})}{\varepsilon^2} = \frac{\sigma^2}{D\varepsilon^2} \xrightarrow{D \rightarrow \infty} 0$$

The CLT gives the form (Gaussian) and the rate ( $1/\sqrt{D}$ ) of convergence.

### Quick Review: Laws of Large Numbers

**WLLN:**  $\bar{X} \xrightarrow{P} \mu$  as  $D \rightarrow \infty$ .

Basis for density estimation.  
**CLT:** convergence rate  $1/\sqrt{D}$ .

---

## 2. Random Projections

**Principle:** Exponentially many quasi-orthogonal directions in high dimensions.

**Johnson-Lindenstrauss Lemma:**

Given  $X \subset \Omega$ , there exists a linear projection  $p : \mathbb{R}^D \rightarrow \mathbb{R}^d$  such that:

$$(1 - \varepsilon)\|x_i - x_j\| \leq \|p(x_i) - p(x_j)\| \leq (1 + \varepsilon)\|x_i - x_j\|$$

with

$$d = O\left(\frac{\log N}{\varepsilon^2}\right)$$

**Important Note:**  $d$  does **NOT** depend on  $D$ !

**Applications:**

- Locally Sensitive Hashing (LSH) for approximate nearest neighbor (ANN) search
- Random bases for signal decomposition and coding

**Caveat:** Convergence rate slower than most concentrations

### Details: Why It Works

**Recall:** Two random vectors  $u$  and  $v$  are likely orthogonal:  $\langle u, v \rangle = \varepsilon$  (small).

**Key Property:** There are exponentially many quasi-orthogonal directions in high-dimensional spaces.

This allows creating random bases (quasi-orthonormal) that preserve essential metric properties while drastically reducing dimension.

### Quick Review: Random Projections

**Johnson-Lindenstrauss:** projection  $\mathbb{R}^D \rightarrow \mathbb{R}^d$  preserving distances with  $d = O(\log N / \varepsilon^2)$  (independent of  $D$ ).

Application: LSH for ANN.

---

## Impact in Data Science

### 1. Sampling

#### Consequences of Sparsity:

- Requires an **exponential** number of samples to cover the space
- Concentration inequalities are **precision indicators** for estimation
- Rejection rate in rejection sampling → close to 100%

#### Details: Rejection Sampling

In practical processes like rejection sampling, the rejection rate becomes exponentially close to 100% due to the empty sphere phenomenon.

If sampling in a particular region (like a sphere inscribed in a cube), the acceptance probability decreases exponentially with  $D$ .

---

### 2. Indexing

#### Exploiting Structure:

- Indexing exploits data structure to anticipate which parts of the space to visit
- If all neighbors are at equal distance, the structure **disappears**

#### Spatial Partitioning Techniques (e.g., kd-trees):

- Attempt to uniquely identify each data point in a cell
- With distance concentration: cell size → full radius
- Search becomes **exhaustive** → complexity  $O(N)$
- Hubness makes some data points answers to almost all queries

#### Details: Search Trees

kd-trees attempt to partition the space so that each data point is uniquely identified within a cell intersection (paths in the tree).

With distance concentration:

- The cell must be as large as the sphere's radius to cover the entire space
  - Pruning conditions become ineffective
  - The tree degenerates into a linear list
-

### 3. Learning

**Machine Learning** = discovering function approximators.

**Problem:** Find  $\phi_\theta \in \mathcal{F}$  approximating  $\phi : X \rightarrow Y$  by minimizing an error.

**Catastrophic Theoretical Bound:**

If  $\mathcal{F}$  is a class of  $c$ -uniformly Lipschitz functions on  $\Omega$ :

$$\sup_{\phi_\theta \in \mathcal{F}} \|\phi_\theta - \phi\|_\infty \geq \frac{c\sqrt{D}N^{-1/D}}{2} \sqrt{\frac{2}{\pi e}} \left(1 + O\left(\frac{\log D}{D}\right)\right)$$

For an error  $c\varepsilon$  with  $\varepsilon > 0$ , the required number of examples is:

$$N \geq \frac{\varepsilon^{-D} D^{D/2}}{(2\pi e)^{D/2}}$$

**Concrete Examples:**

- $D = 3, \varepsilon = 0.1 \rightarrow N \geq 74$
- $D = 10, \varepsilon = 0.1 \rightarrow N \geq 687 \times 10^6$
- $D = 15, \varepsilon = 0.1 \rightarrow N \geq 377 \times 10^{12}$
- $D = 10, \varepsilon = 0.01 \rightarrow N \geq 6.9 \times 10^{18}$

**Realistic Case:**  $D \sim 100, \varepsilon \sim 10^{-4} \rightarrow N \sim 10^{400}$  **IMPOSSIBLE**

#### Quick Review: Importance in Data Science

**Sampling:** Exponential need, rejection rate  $\rightarrow 100\%$ .

**Indexing:** Structure disappears, kd-trees  $\rightarrow O(N)$ , hubness.

**Learning:**  $N \geq \varepsilon^{-D} D^{D/2} / (2\pi e)^{D/2} \rightarrow$  impossible for large  $D$ .

---

### 4. Why Does Data Science Still Work?

**Assumptions in the Above Analyses:**

- **Uniform** (or isotropic) data distribution
- **Independent** features

**These assumptions are FALSE in practice!**

**Reality:**

- Real data **do NOT** populate the space uniformly
- Features show **correlations**
- Data occupy **subspaces** of lower intrinsic dimensionality

#### Practical Solutions:

- Decorrelate features (linear latent space)
- Discover subspaces (clustering, nonlinear dimensionality reduction)
- PCA, autoencoders, neural networks
- Exploit underlying geometric structure

**Key Insight:** If features are informative, they are not independent, and data concentrate on manifolds of much lower dimension than  $D$ .

#### Quick Review: Why It Works

Analyses assume **uniform** distribution and **independent** features.

**Reality:** Data lie on **subspaces** of low intrinsic dimensionality.

**Solutions:** Decorrelation, clustering, dimensionality reduction.

---

## Global Summary

High-dimensional geometry ( $D \gtrsim 10$ ) is **different** from familiar 3D geometry.

### Positive Aspects (Blessing)

- **Concentration inequalities:** Basis of statistical estimation
- **Laws of Large Numbers:** Enable density estimation
- **Random projections:** Exploit quasi-orthogonality

### Negative Aspects (Curse)

- Data highly **sparse** (exponential sample requirement)
- **Concentration of distances and angles:**
  - All neighbors equidistant
  - Everything appears orthogonal
- Density concentrated on **narrow radii**  $\rightarrow$  biased sampling
- **Hubs:** Part of all classes/responses

## Unique Solution

Since this is by construction, the only solution is to operate in **low-dimensional spaces**:

- Dimensionality reduction
- Subspace discovery
- Feature decorrelation

**Key Message:** Data Science works because real data **are not uniform** and live on **low-dimensional submanifolds**.

### Quick Review: Summary

**High dimension** ( $D \gtrsim 10$ ) 3D geometry.

: Concentration inequalities, WLLN, random projections.

: Data sparsity, distance/angle concentration, hubs.

**Solution:** Operate in **low dimension** (PCA, autoencoders, structure exploitation).