

Table of contents

The Essentials	2
Introduction and Motivation	2
The Fundamental Problem	2
Geometric Intuition: The Point Cloud	2
Maximizing Variance	2
Mathematical Formalization	2
Complete Decomposition	3
Minimizing Error	3
Variance-Error Equivalence	4
Interpretations	4
Dimensionality Reduction	4
Structure of the Latent Space	4
Latent Coordinates	5
Low-Rank Approximation	5
Eckart-Young Theorem	5
Practical Truncation	6
Geometry and Reconstruction Quality	6
Relative Contribution to Variance	6
Projection Ratio	6
Normalization and Scaling	7
Practical Application of PCA	7
Complete PCA Algorithm	7
Choosing the Number of Components K	7
Linear Autoencoders	8
Structure of a Linear Autoencoder	8
Extensions of PCA	9
Probabilistic PCA	9
Kernel PCA	9
Local PCA	10
What PCA Tells You (and What It Doesn't)	10
What It Tells You	10
What It Doesn't Tell You	10
When to Use PCA?	10
Suitable For	10
Not Suitable For	11

The Essentials

Introduction and Motivation

PCA (Principal Component Analysis) is a **linear latent decomposition** method that identifies the **underlying latent factors** structuring data.

The Fundamental Problem

Base representations of data are not always optimal. PCA seeks to find a new representation that:

- Reveals the **latent factors** (hidden) governing the process
- Uses **statistical correlation** to identify these factors
- Enables **simplification** of data through justified decimation

Key Point: PCA belongs to the family of **linear latent models**

Geometric Intuition: The Point Cloud

Imagine a 3D point cloud (like midges in a room). If you had to squash them onto a wall:

- Which wall would you choose so that the midges are **most spread out**?
- That wall is the **principal plane**
- The line in this wall where they are most aligned is the **first component**

Essential Idea: PCA finds the **skeletal structure** of the data, like a radiologist seeing bones through skin.

Maximizing Variance

Mathematical Formalization

Idea: We seek a line passing through the center where the projections are **as dispersed as possible**.

Let X be the dataset, $\{u_i\}_{i \in \llbracket D \rrbracket}$ a new **orthonormal basis** of $\mathbb{R}^D$.

The **first principal component** u_1 is chosen such that the **variance of the projected data** on u_1 is maximized.

i Mathematical Details: Optimization Problem

The empirical mean and projected variance are:

$$x = \frac{1}{N} \sum_{i=1}^N x_i$$
$$v_{u_1} = \frac{1}{N} \sum_{i=1}^N (u_1^T x_i - u_1^T x)^2 = u_1^T \Sigma u_1$$

where $\Sigma = \frac{1}{N} \sum_{i=1}^N (x_i - x)(x_i - x)^T$ is the **covariance matrix** of the data.
The problem becomes:

$$u_1 = \arg \max_{u^T u = 1} u^T \Sigma u$$

Using **Lagrange multipliers**:

$$J(u) = u^T \Sigma u + \lambda(1 - u^T u)$$

The optimality conditions give:

$$\Sigma u_1 = \lambda_1 u_1 \quad \Rightarrow \quad v_{u_1} = \lambda_1$$

Conclusion: (u_1, λ_1) is an **eigenpair** of the covariance matrix Σ

Complete Decomposition

Continuing this process, we obtain the set of **principal components** as the eigenpairs $\{(u_i, \lambda_i)\}_{i \in \llbracket D \rrbracket}$ of Σ .

Matrix Decomposition:

$$\Lambda = U^T \Sigma U$$

where U contains the eigenvectors in columns and Λ is diagonal with the eigenvalues.

Minimizing Error

Idea: Represent each point by its projection on a line with **minimal reconstruction error** (perpendicular distance).

Variance-Error Equivalence

The variance v_{u_1} is expressed as a sum of squares:

$$v_{u_1} = \frac{1}{N} \sum_{i=1}^N (u_1^T(x_i - x))^2$$

i Mathematical Details: Pythagorean Theorem

To maximize v_{u_1} , the terms $u_1^T(x_i - x)$ must be collectively maximized. Since $(x_i - x)$ is fixed, this is equivalent to **minimizing the distance to the projection axis** (approximation error).

Alternative Formulation:

$$u_1 = \arg \min_{u^T u = 1} \sum_{i=1}^N \|(x_i - x) - [u^T(x_i - x)]u\|_2^2$$

This formulation also gives: $\Sigma u_1 = \lambda_1 u_1$

Fundamental Equivalence:

Maximizing projection variance Minimizing approximation error

Interpretations

- **Regression:** A principal component is a **quadratic regression** on the data
- **Physics:** A principal component is a **subspace of minimal inertia** with respect to X

Dimensionality Reduction

Structure of the Latent Space

The latent space with basis $\{u_1, \dots, u_D\}$ has the following properties:

- **Eigenvalue Order:** By construction: $\lambda_1 > \lambda_2 > \dots > \lambda_D$
- **Total Variance:** $\lambda = \sum_{d=1}^D \lambda_d = \text{Tr}(\Sigma)$
- **Decorrelation:** $v_{u_d} = \lambda_d$ and $u_d^T u_{d'} = 0$ for $d \neq d'$

Latent Coordinates

The basis $\{u_1, \dots, u_D\}$ induces **latent coordinates** y_i for the data:

$$y_i = U^T(x_i - x)$$

The PCA transformation is **linear** and consists of three steps:

- **Centering** on x
- **Rotation** using matrix U
- **Scaling** (whitening) using $\Lambda^{1/2}$

Important Limitation: PCA assumes an approximately **Gaussian distribution** and will not be relevant for non-Gaussian data (e.g., clustered data).

Low-Rank Approximation

Eckart-Young Theorem

If $\Sigma = U\Lambda U^T$ and for $K < D$, we define:

$$\Sigma_K = \sum_{d=1}^K \lambda_d u_d u_d^T$$

Then Σ_K is the **closest rank- K matrix** to Σ in the Frobenius norm.

i Mathematical Details: Approximation Error

The theorem states that:

$$\arg \min_{\text{rank}(S)=K} \|\Sigma - S\|_F^2 = \Sigma_K$$

where $\|\cdot\|_F$ is the **Frobenius norm**.

The approximation error is:

$$\|\Sigma - \Sigma_K\|_F^2 = \sum_{d=K+1}^D \lambda_d$$

This error corresponds exactly to the variance lost during reduction.

Practical Truncation

Truncation of U and Λ :

$$\Sigma_K = U_K \Lambda_K U_K^T$$

Projection and Reconstruction:

- **Projection** (encoding): $\tilde{y}_i = U_K^T x_i \in \mathbb{R}^K$
- **Reconstruction** (decoding): $\tilde{x}_i = U_K U_K^T x_i + x$

Important: We go from D dimensions to K dimensions with $K < D$

Geometry and Reconstruction Quality

Relative Contribution to Variance

The contribution of each dimension:

$$c_d = \frac{\lambda_d}{\sum_k \lambda_k}$$

This metric depends on the **spectrum distribution** of eigenvalues.

Projection Ratio

The projection ratio of a data point x_i on a latent factor:

$$\rho_d(x_i) = \frac{\langle u_d, x_i \rangle^2}{\|x_i\|^2} = \cos^2(\angle(u_d, x_i))$$

The closer $\rho_d(x_i)$ is to 1, the more x_i lies on the axis u_d .

Normalization and Scaling

PCA is more effective if **all scales are similar**.

Solution: Normalization

We create the **scaling matrix** $S = \text{diag}[\sigma_1^2, \dots, \sigma_D^2]$ and use the **Mahalanobis metric**:

$$d_S^2(x, y) = (x - y)^T S^{-1} (x - y)$$

Practical Tip: Normalizing data before PCA often improves results.

Practical Application of PCA

Complete PCA Algorithm

Given dataset $X \in \Omega$:

- Compute the empirical mean $\bar{x} = \frac{1}{N} \sum_{i=1}^N x_i$
- Center the data: $x_i \leftarrow (x_i - \bar{x})$ and form the centered matrix X
- Compute the covariance matrix: $\Sigma = \frac{1}{N} X X^T$
- Decompose into eigenvalues: $\Sigma = U \Lambda U^T$

So far: **exact** transformation

- Select the number of components K
- Define U_K , Λ_K and compute $\{\tilde{y}_i\}_{i \in \llbracket N \rrbracket}$ and/or $\{\tilde{x}_i\}_{i \in \llbracket N \rrbracket}$

Choosing the Number of Components K

Three Main Strategies:

Strategy 1: Target Dimension

$$K = d$$

where d is the desired dimension.

Strategy 2: Variance Threshold

Require that the latent space retains a proportion τ of the total variance:

$$\frac{\sum_{d=1}^K \lambda_d}{\sum_{d=1}^D \lambda_d} > \tau$$

Example: $\tau = 0.95$ means retaining 95% of the variance.

Strategy 3: Reconstruction Error

Train K such that $L(X) \leq \epsilon$, where for example:

$$L(X) = \sum_i \|x_i - \tilde{x}_i\|^2$$

Practical Advice: Visualizing the eigenvalue spectrum (scree plot) helps identify a natural “elbow”.

Linear Autoencoders

Structure of a Linear Autoencoder

A linear autoencoder has:

- W_{enc} and W_{dec} : weight matrices (encoder and decoder)
- $z(x_i) = W_{\text{enc}} x_i$: encoding
- $\tilde{x}_i = W_{\text{dec}} z(x_i)$: decoding

Optimization Problem:

$$\theta^* = \arg \min_{\text{rank}(W)=K} \|X - WZ\|_F^2$$

i Mathematical Details: Relationship Between PCA and Autoencoder

Using the Eckart-Young theorem with SVD decomposition $X = U\Psi V^T$, the solution is:

$$W_{\text{dec}} Z = U_K \Psi_K V_K^T$$

Setting $W_{\text{dec}} = U_K \Psi_K$, then:

$$W_{\text{enc}} = \Psi_K^{-1} U_K^T$$

For **centered** data:

- **PCA:** $\Sigma = \frac{1}{N}XX^T = U\Lambda U^T$ thus $\tilde{x}_i = U_K U_K^T x_i$
- **AE:** $X = U\Psi V^T$ thus $\Sigma = \frac{1}{N}U\Psi^2 U^T$ and $\tilde{x}_i = U_K U_K^T x_i$

Relationship between eigenvalues:

$$\Lambda = \frac{1}{N}\Psi^2$$

Fundamental Conclusion: A linear autoencoder performs PCA if the data is **centered** and **normalized**.

Note: W_{enc} is **not the inverse** of W_{dec} if $K < D$.

Extensions of PCA

Probabilistic PCA

Reformulation with explicit latent distribution:

- Latent distribution: $\mathbb{P}(z) = \mathcal{N}(0, I_{dK})$
- Conditional model: $\mathbb{P}(x|z) = \mathcal{N}(Wz + \mu, \sigma^2 I_{dD})$

Advantages:

- Accommodation of **measurement noise** (variance σ^2)
- **Generative mode:** possibility to generate new data
- Estimation by **Maximum Likelihood**
- **EM algorithm** to save computations

Kernel PCA

Use of a **nonlinear mapping** $\phi(x_i)$ via a kernel function:

$$k(x_i, x_j) = \phi(x_i)^T \phi(x_j)$$

Allows performing PCA in a **more favorable space** (possibly infinite-dimensional) without explicitly computing ϕ .

Local PCA

Performs PCA on **local neighborhoods** of the data to consider only the **local intrinsic dimensionality**.

Useful for data with complex local structure (nonlinear manifolds).

What PCA Tells You (and What It Doesn't)

What It Tells You

- The **linear structure** of the data
- The **principal directions** of variation
- The **intrinsic dimension** (how many dimensions are truly useful)
- Which **variables** contribute to each axis

What It Doesn't Tell You

- **Not clustering**: It doesn't find groups
 - **Linear only**: It misses curved relationships
 - **Sensitive to outliers**: An extreme point can shift everything
-

When to Use PCA?

Suitable For

- Data approximately following a Gaussian distribution
- Dimensionality reduction with variance preservation
- Visualization of high-dimensional data
- Denoising data corrupted by Gaussian noise
- Data compression

Not Suitable For

- Strongly clustered data (non-Gaussian)
 - Data with strongly nonlinear relationships
 - Cases where variance is not a good information criterion
-

Condensed Cheat Sheet

In 30 seconds:

- PCA = smart basis change
- New axes = directions of maximum variance
- Eigenvalues = importance of each axis
- Always center first!
- Choose K with the scree plot
- Useful for: visualization, dimensionality reduction, denoising
- Not useful for: clustering, nonlinear relationships

Key Mathematical Concepts to Master

To understand PCA well:

- **Linear transformation** and **orthogonal matrix**
- **Coordinates** and **matrix rank**
- **Mean** and **variance**
- **Orthogonal projection**
- **Eigendecomposition**
- **Matrix trace**
- **Frobenius norm**
- **Lagrange multipliers**