

Table of contents

Detailed Course	2
Introduction and Motivation	2
General Context	2
Supervised Learning: Foundations	2
General Definition	2
Examples of Learners	3
Principle of LDA: Discriminant Direction	3
Problem Statement	3
Discriminant Direction with 2 Classes	3
Between-Class Criterion	4
Principle of Maximum Separation	4
Mathematical Formulation	4
Within-Class Criterion	5
Principle of Minimum Variance	5
Within-Class Variance	5
Fisher's Criterion	5
Generalization to M Classes	6
Generalized Between-Class Criterion	6
Generalized Within-Class Criterion	6
Combined Fisher's Criterion	7
Combining Both Criteria	7
Solution by Differentiation	7
Discriminant Subspaces	7
Multiple Discriminant Dimensions	7
Dimensional Limitation	8
Special Case: $M = 2$	8
Practical Example: Character Recognition	8
Dataset Description	8
Result of LDA Projection	9
Inference and Decision Making	9
Conditional Modeling	9
Probabilistic Approach	9
Gaussian Approximation	10
Class-Specific Gaussian Model	10
Maximum Likelihood	10
Decision Rule	10
Optimality of LDA	11
Optimality Conditions	11
Limitations	11

Synthesis of Key Concepts	11
Essential Characteristics of LDA	11
Comparison: LDA vs PCA	12
Fundamental Differences	12
Complementarity	12
Typical Exam Questions	13
Conceptual Questions	13
Technical Questions	13
Mathematical Proofs	14
Mathematical Concepts to Master	14

Detailed Course

Introduction and Motivation

General Context

LDA (Linear Discriminant Analysis) is a **supervised latent analysis** method that finds the optimal latent factors for discrimination between classes.

Objectives of LDA:

- Introduce a **supervised** definition of latent factors
- Propose a **conditional modeling** of data
- Perform **classification** of unlabeled items

Warning: This LDA should not be confused with **Latent Dirichlet Allocation** (used in NLP)

Supervised Learning: Foundations

General Definition

Given data $X \subset \Omega \subseteq \mathbb{R}^D$ associated with categories (class labels) $Y = [[M]] \subset \mathbb{N}$.

Supervised learning consists of finding (learning) the parameters θ of a learner (function) ϕ_θ to minimize the loss $L_{X \times Y}(\theta)$ incurred when predicting class labels.

Prediction Function:

$$\begin{aligned}\phi_\theta &: \Omega \rightarrow Y \\ x_i &\mapsto \phi_\theta(x_i) = \tilde{y}_i\end{aligned}$$

Examples of Learners

Logistic Regression:

ϕ_θ is a logistic regression with parameters θ

Neural Network:

ϕ_θ is a neural network with weights θ

Example Loss Function:

$$L_{X \times Y}(\theta) = \sum_X \|\phi_\theta(x_i) - y_i\|^2$$

Principle of LDA: Discriminant Direction


Problem Statement

Fundamental Question: Given multivariate data X associated with labels Y , we seek a model for this data.

Discriminant Direction with 2 Classes

We seek a **direction** $a \in \mathbb{R}^D$ such that the projection of the data onto this direction $\hat{X} = \text{Proj}_a(X)$ shows **optimal properties for discrimination** between classes.

Geometric Visualization:



Key Principle: Find the direction a that best separates the classes after projection

Between-Class Criterion

Principle of Maximum Separation

If the projections of the classes are **well separated**, then they can be easily discriminated.

Objective: Seek the direction a where between-class discrimination is **maximized**

Mathematical Formulation

Geometric Intuition:

The direction a should be **collinear** with the line $\mu_1 - \mu_2$ connecting the class centroids, because:

$$\|\hat{\mu}_1 - \hat{\mu}_2\|^2 = \left\| \frac{a^T \mu_1}{\|a\|^2} a - \frac{a^T \mu_2}{\|a\|^2} a \right\|^2 = \frac{1}{\|a\|^2} (a^T (\mu_1 - \mu_2))^2$$

This expression is **maximized** when $a = \lambda(\mu_1 - \mu_2)$

Important Point: The original data ($x_i \in X$) does not change. Only the **projected data** ($\hat{x}_i \in \hat{X}$) and subsequent calculations change as a evolves.

Property: The centroid of the projected data is the projection of the centroid of the original data

Within-Class Criterion

Principle of Minimum Variance

We also seek to **minimize the variance** of each projected class individually.

Objective: Reduce dispersion within each class after projection

Within-Class Variance

For each class k , the normalized variance of the projection is:

$$\hat{\sigma}_k^2 = \frac{1}{N_k} \sum_{y_i=k} (\hat{x}_i - \hat{\mu}_k)^T (\hat{x}_i - \hat{\mu}_k)$$

Fisher's Criterion

Fisher's criterion maximizes the ratio:

$$\arg \max_a \frac{\|\hat{\mu}_1 - \hat{\mu}_2\|^2}{\hat{\sigma}_1^2 + \hat{\sigma}_2^2}$$

Interpretation:

- **Numerator:** Distance between projected means (between-class) $\rightarrow$ to **maximize**
 - **Denominator:** Sum of within-class variances $\rightarrow$ to **minimize**
-

Generalization to M Classes

Generalized Between-Class Criterion

Definitions:

Let $B = [\mu_1 - \mu, \dots, \mu_M - \mu]$ be the matrix of centered class centroids, where:

- $\mu = \frac{1}{N} \sum_i x_i$ is the global mean
- $N = \sum_k N_k$ is the total number of points

Between-class covariance matrix:

$$\Sigma_b = \frac{1}{M} B B^T$$

Criterion to maximize:

$$\frac{1}{M} \sum_k (\hat{\mu}_k - \hat{\mu})^T (\hat{\mu}_k - \hat{\mu}) = \frac{1}{\|a\|^2} a^T \Sigma_b a$$

Generalized Within-Class Criterion

Definitions:

Let $A_k = [x_1 - \mu_k, \dots, x_{N_k} - \mu_k]$ with $y_i = k$, be the matrix of centered data for class k .

Within-class covariance matrix:

- For class k : $\frac{1}{N_k} A_k A_k^T$
- **Sum of within-class matrices:** $\Sigma_w = \sum_k \frac{1}{N_k} A_k A_k^T$

Criterion to minimize:

$$\sum_k \frac{1}{N_k} \sum_{y_i=k} (\hat{x}_i - \hat{\mu}_k)^T (\hat{x}_i - \hat{\mu}_k) = \frac{1}{a^T a} a^T \Sigma_w a$$

Development:

$$\sum_k \frac{1}{N_k} \sum_{y_i=k} \left[\frac{a^T (x_i - \mu_k)}{\|a\|^2} a \right]^T \left[\frac{a^T (x_i - \mu_k)}{\|a\|^2} a \right] = \frac{1}{a^T a} a^T \Sigma_w a$$

Combined Fisher's Criterion

Combining Both Criteria

Generalized Fisher's criterion combines the between-class and within-class criteria:

$$\arg \max_a L(a) = \arg \max_a \frac{a^T \Sigma_b a}{a^T \Sigma_w a}$$

Solution by Differentiation

To find the maximum, we compute the derivative:

$$\frac{\partial L(a)}{\partial a} = \frac{\Sigma_b a (a^T \Sigma_w a) - \Sigma_w a (a^T \Sigma_b a)}{(a^T \Sigma_w a)^2} = 0$$

Resulting Equation:

$$\Sigma_b a = L(a) \Sigma_w a$$

Solution: a is a solution of the **generalized eigenvalue problem**

The optimal direction a is the **first eigenvector** of $\Sigma_w^{-1} \Sigma_b$ (with eigenvalue $\lambda_1 = L(a)$)

Discriminant Subspaces

Multiple Discriminant Dimensions

Ordered Eigenvectors:

The eigenvectors corresponding to the **largest eigenvalues** λ_i are the most discriminative dimensions:

$$a_1, a_2, \dots, a_p \quad \text{with} \quad \lambda_1 > \lambda_2 > \dots > \lambda_p$$

Dimensional Limitation

Fundamental Property: M classes can be discriminated in a subspace of dimension **at most** $(M - 1)$.

Consequence: There are only $M - 1$ **non-zero** eigenvalues

Explanation: With M classes, we have M class centroids. Once globally centered, these centroids are constrained to lie in a space of dimension $M - 1$.

Special Case: $M = 2$

For two classes, we can show that:

$$BB^T = (\mu_1 - \mu)(\mu_1 - \mu)^T + (\mu_2 - \mu)(\mu_2 - \mu)^T \propto (\mu_1 - \mu_2)(\mu_1 - \mu_2)^T$$

Implications:

- $BB^T a$ is a vector in the direction of $(\mu_1 - \mu_2)$
 - The solution is $a = \lambda \Sigma_w^{-1}(\mu_1 - \mu_2)$
-

Practical Example: Character Recognition

Dataset Description

Partial MNIST Dataset:

- 7291 images of size 16×16 (8 bits)
- Digits from 0 to 9
- Representation: $\{x_i, y_i\}$ with $x_i \in \mathbb{R}^{256}$ and $y_i \in [[10]]$, $i \in [[7291]]$

Result of LDA Projection

LDA finds the **optimal subspace** for **linearly** separating the data according to their labels y_i .

Advantages of LDA Projection:

- Dimensionality reduction (from 256 to a few dimensions)
 - Visualization in 2D or 3D possible
 - Class separation maximized
-

Inference and Decision Making

Conditional Modeling

Classification Problem:

For a new data point j :

- x_j is **known**
- y_j is **unknown**

Question: To which class k does point j belong?

Probabilistic Approach

Declaration of a Categorical Random Variable:

We declare the class random variable C

Prediction by Bayes' Rule:

$$\mathbb{P}(C = k|x_j) = \frac{\mathbb{P}(x_j|C = k)\mathbb{P}(C = k)}{\mathbb{P}(x_j)}$$

Components:

- $\mathbb{P}(x_j|C = k)$: **likelihood**
 - $\mathbb{P}(C = k)$: **prior** (a priori probability)
 - $\mathbb{P}(x_j)$: **evidence** (normalization)
-

Gaussian Approximation

Class-Specific Gaussian Model

Distribution Assumption:

Each class is modeled by a normal distribution:

$$\mathbb{P}(x|C = k) \sim \mathcal{N}(\mu_k, \Sigma_k)$$

Prior Choice:

Uniform prior: $\mathbb{P}(C = k) = \frac{1}{M}$ (all classes equally probable)

Simplification:

The evidence $\mathbb{P}(x_j)$ is **ignored** as it is constant for all classes (normalization)

Maximum Likelihood

Approximation:

$$\mathbb{P}(x|C = k) \propto \exp\left(-\frac{1}{2}(x - \mu_k)^T \Sigma_k^{-1} (x - \mu_k)\right)$$

Decision Rule

Classification by Maximum A Posteriori:

$$\delta(x) = \arg \max_k \mathbb{P}(x|C = k)$$

Interpretation: We assign x to the class k that maximizes the likelihood (or posterior probability)

Optimality of LDA

Optimality Conditions

LDA is **optimal** when the M classes each follow a **Gaussian distribution**

Justification:

The discrimination criteria are based on covariance matrices Σ_w and Σ_b , which fully characterize Gaussian distributions.

Limitations

Linear Discriminant Analysis: does **not account** for **nonlinear** relationships between variables

Linear discrimination only

Alternative: For nonlinear relationships, use **nonlinear classification** methods:

- Kernel LDA
 - Quadratic Discriminant Analysis (QDA)
 - Support Vector Machines (SVM)
 - Neural networks
-

Synthesis of Key Concepts

Essential Characteristics of LDA

Nature of the Method:

- LDA is a **supervised** technique
- Uses **class labels** to guide dimensionality reduction

Objective:

- Find the **optimal linear latent factors** that explain linear discrimination between classes

Fisher's Criterion:

- Combines **within-class** (within-scatter) and **between-class** (between-scatter) dispersion

- Maximizes the ratio $\frac{a^T \Sigma_b a}{a^T \Sigma_w a}$

Solution:

- Solved by finding the eigenvalues of the matrix $\Sigma_w^{-1} \Sigma_b$
- The eigenvectors form the discriminant axes

Applications:

- **Visualization** of data considering labels
- **Supervised dimensionality reduction**
- **Classification** via conditional modeling (each class follows a normal distribution)

Comparison: LDA vs PCA

Fundamental Differences

PCA (Unsupervised):

- Maximizes **global variance**
- Ignores class labels
- Finds directions of maximum variance
- Based on Σ (total covariance)

LDA (Supervised):

- Maximizes **separation between classes**
- Uses class labels
- Finds discriminant directions
- Based on $\Sigma_w^{-1} \Sigma_b$ (ratio of covariances)

Complementarity

In Practice: PCA and LDA can be combined

1. PCA for initial dimensionality reduction (eliminate noise)
2. LDA to maximize discrimination

Typical Exam Questions

Conceptual Questions

Supervised Learning Setup:

- Explain the context of supervised learning
- Difference between supervised and unsupervised learning

Linear Nature of LDA:

- Why is LDA a linear method?
- What are its limitations?

Discriminant Direction:

- What characterizes a discriminant direction?
- What does the vector a represent?

Technical Questions

Fisher's Criteria:

- What is Fisher's criterion?
- How does it combine between-class and within-class criteria?

Between-Class Criterion:

- Principle of the between-class criterion
- How to compute the between-class covariance matrix Σ_b ?

Within-Class Criterion:

- Principle of the within-class criterion
- How to compute the within-class covariance matrix Σ_w ?

Mathematical Proofs

Derivation of the Maximum:

- Justify and explain the derivation of the maximum for a
- Show that a satisfies $\Sigma_b a = \lambda \Sigma_w a$

Multidimensional Situation:

- Describe the multidimensional situation ($D > 2$)
- Explain why there are at most $M - 1$ discriminant dimensions

Inference:

- How to perform inference with LDA?
 - Provide an example of conditional modeling using LDA
-

Mathematical Concepts to Master

To understand LDA well, you need to master:

Linear Algebra:

- Eigenvectors and eigenvalues
- Generalized eigenvalue problem
- Covariance matrices
- Orthogonal projection

Statistics:

- Mean and variance
- Covariance and correlation
- Multivariate normal distribution
- Bayes' rule

Optimization:

- Differentiation of matrix functions
- Constrained optimization
- Rayleigh quotient method