

Table of contents

Detailed Course	2
Introduction and Motivation	2
General Context	2
Recap: Models Used So Far	2
New Approach: Density Modeling	3
Density Estimation Methods	3
Mixture Models	3
Definition	3
Probabilistic Interpretation	4
Maximum Likelihood Estimation (MLE)	4
Gaussian Mixture Model (GMM)	5
Choice of Gaussian Family	5
Advantages of GMM	6
Simple Case: 1 Component	6
Complex Case: 2 Components	7
Latent Variables and EM Strategy	7
E-step (Expectation)	8
Hard vs. Soft Assignment	9
M-step (Maximization)	9
EM Algorithm Convergence	10
EM Algorithm Limitations	11
Generalization to c Components	12
Complete EM Algorithm for c Components	12
Step 1: Initialization	12
Step 2: E-step	13
Step 3: M-step	13
Step 4: Iteration	13
Practical Modeling	14
Covariance Matrix Form	14
Preprocessing: Dimensionality Reduction with PCA	15
Selecting the Number of Components	16
Synthesis of Key Concepts	17
Gaussian Mixture Models (GMM)	17
EM Algorithm	17
Relationship Between GMM and K-means	18
Conceptual Comparison	18
GMM as a Generalization of K-means	18
Typical Exam Questions	18
Conceptual Questions	18
Technical Questions	19

Mathematical Developments	19
Mathematical Concepts to Master	19
Probability and Statistics	19
Optimization	20

Detailed Course

Introduction and Motivation

General Context

The **Expectation-Maximization (EM) algorithm** is a method for **estimating latent factors** in an **unsupervised** context of conditional modeling.

Course Objectives:

- Understand **unsupervised conditional modeling** as latent factor estimation
- Develop the example of **density estimation using Gaussian Mixture Models (GMM)**
- Practice **Maximum Likelihood Estimation (MLE)**
- Understand the **alternative principle** of Expectation-Maximization for latent factor discovery

Recap: Models Used So Far

Component Models (PCA):

- Criterion based on the **variance** of centered data
- Implicit distribution: **normal**

Discriminant Models (LDA):

- Variance of projected data for within-class and between-class models

Common Point: The distribution is **fixed** (essentially normal) and we seek its parameters θ (e.g., $\theta = [\mu, \Sigma]$)

New Approach: Density Modeling

Alternative: Seek a model for the **data density**

Let $f_\theta(x) : \Omega \rightarrow \mathbb{R}$ be the **data density**

Density Estimation Methods

Overview of Methods

Non-parametric Methods:

- **k-Nearest Neighbors (kNN):** Estimation based on the k nearest neighbors
- **Parzen Windows:** Kernel-based estimation
- **RBF Networks:** Radial basis functions
- **Histograms:** Space discretization

Parametric Methods:

- **Mixture Models:** The focus of this course
-

Mixture Models

Definition

General Formulation

Definition:

The density $f(x)$ is generated by c “basis” functions ϕ (components), from a family $\mathcal{F}$:

$$f(x) = \sum_{j=1}^c \pi_j \phi(x, \theta_j)$$

Model Components:

- $\pi_j \in \mathbb{R}$: **mixture parameters**
- $\phi(x, \theta_j) \in \mathcal{F}$: **basis functions** controlled by parameters θ_j

Model Assumptions

Three Fundamental Assumptions:

1. The **number of components** $c \in \mathbb{N}^*$ is **known**
 2. The **function family** $\mathcal{F}$ is **known**
 3. The **labels** (class labels) are **unknown** (unsupervised context)
-

Probabilistic Interpretation

Probabilistic Reading of the Mixture

The density $f(x)$ represents a **random process** where x is drawn from a set of states ω_j with prior probability $\mathbb{P}(\omega_j)$.

Probabilistic Formulation:

$$f(x) = \sum_{j=1}^c \pi_j \phi(x, \theta_j) = \mathbb{P}(x|\theta) = \sum_{j=1}^c \mathbb{P}(\omega_j) \mathbb{P}(x|\omega_j, \theta_j)$$

By Identification:

- $\pi_j = \mathbb{P}(\omega_j)$ with the constraint $\sum_j \pi_j = 1$
 - $\phi(x, \theta_j) = \mathbb{P}(x|\omega_j, \theta_j)$ is the **conditional probability** that x is generated by state ω_j
-

Maximum Likelihood Estimation (MLE)

Likelihood Recap

Data:

Let $X = \{x_1, \dots, x_N\}$ be **unlabeled** samples generated by the mixture $f(x) = \mathbb{P}(x|\theta)$

Parameters to Infer:

$$\theta = [\pi_j, \theta_j]_{j \in [[c]]}$$

Likelihood Function

Likelihood:

Assuming the data are **independent and identically distributed (i.i.d.)**:

$$\mathbb{P}(X|\theta) \stackrel{\text{i.i.d.}}{=} \prod_{i=1}^N \mathbb{P}(x_i|\theta)$$

Likelihood Estimator:

$$\hat{\theta} = \arg \max_{\theta} \mathbb{P}(X|\theta)$$

Log-Likelihood

Logarithmic Transformation:

$$LL(\theta, X) = \sum_{i=1}^N \log \mathbb{P}(x_i|\theta) = \sum_{i=1}^N \log \left[\sum_{j=1}^c \pi_j \phi(x_i, \theta_j) \right]$$

Maximum Log-Likelihood Estimator (MLE):

$$\hat{\theta} = \arg \max_{\theta} LL(\theta, X)$$

Advantage: Transforms the product into a sum, numerically more stable

Gaussian Mixture Model (GMM)

Choice of Gaussian Family

Since ϕ must be a **probability density function**, choosing the **normal density family** as a basis seems reasonable (an **agnostic** choice):

$$\phi \in \mathcal{F} = \{\mathcal{N}(\mu, \Sigma)\}$$

Gaussian Mixture Model:

$$f(x) = \sum_{j=1}^c \pi_j \mathcal{N}(x|\mu_j, \Sigma_j)$$

where $\phi_j(x) := \phi(x, \theta_j) = \mathcal{N}(x|\mu_j, \Sigma_j)$

Advantages of GMM

Universality: Can approximate **any density** (universal approximation theorem)

Linearity: Allows a **linear system** of parameters when maximizing the log-likelihood

Simple Case: 1 Component

Direct Estimation

Parameters: $\theta = [\mu, \Sigma]$

Optimization Problem:

$$\hat{\theta} = \arg \max_{\theta} \sum_i \log e^{-(x_i - \mu)^T \Sigma^{-1} (x_i - \mu)} \Leftrightarrow \hat{\theta} = \arg \min_{\theta} \sum_i (x_i - \mu)^T \Sigma^{-1} (x_i - \mu)$$

Analytical Solutions:

$$\hat{\mu} = \frac{1}{N} \sum_i x_i$$

$$\hat{\Sigma} = \frac{1}{N} \sum_i (x_i - \hat{\mu})(x_i - \hat{\mu})^T$$

Note: Solutions identical to standard empirical estimation of mean and covariance

Complex Case: 2 Components

Problem Statement

Parameters:

$$\theta = [\pi_1, \theta_1, \pi_2, \theta_2] = [\pi_1, [\mu_1, \Sigma_1], \pi_2, [\mu_2, \Sigma_2]]$$

with the constraint $\pi_1 + \pi_2 = 1$

Log-Likelihood:

$$LL(\theta, X) = \sum_{i=1}^N \log[\pi_1 \phi(x_i, \theta_1) + \pi_2 \phi(x_i, \theta_2)]$$

Difficulty: Maximization is challenging due to the **sum inside the logarithm!**

Solution: EM Algorithm

Solution: Iterative algorithm in **2 steps** (E-M) to maximize LL

→ **Expectation-Maximization (EM) algorithm**

Latent Variables and EM Strategy

Introduction of Latent Variables

Fundamental Problem:

What is missing here is the **assignment** of x_i to one of the 2 components ϕ_j

Key Insight:

If we **knew** this assignment, we would treat the problem as **two instances of the 1-component case**

Binary Assignment Variables

We introduce **latent (unobserved) variables**: the binary assignment $\delta_{ij} \in \{\text{true}, \text{false}\}$ of each x_i to component ϕ_j :

- $x_i \sim \phi_1 \Rightarrow \delta_{i1} = \text{true}, \delta_{i2} = \text{false}$
- $x_i \sim \phi_2 \Rightarrow \delta_{i1} = \text{false}, \delta_{i2} = \text{true}$

But we can only **estimate** $\delta_{ij} \rightarrow$ **EM strategy**

E-step (Expectation)

E-step Principle

Parameters:

$$\theta = [\pi_1, \theta_1, \pi_2, \theta_2] = [\pi_1, [\mu_1, \Sigma_1], \pi_2, [\mu_2, \Sigma_2]]$$

Initialization:

We assume an initial value is known:

$$\theta^0 = [\pi_1^0, \theta_1^0, \pi_2^0, \theta_2^0]$$

Responsibility Calculation

We can infer the **statistical contribution** γ_{ij} of each data point x_i to each component ϕ_j (parameterized by θ_j^0):

$$\gamma_{ij}(\theta^0) = \mathbb{P}[\delta_{ij} = \text{true} | \theta^0, X]$$

Bayes' Formula:

$$\gamma_{i1}(\theta^0) = \frac{\pi_1^0 \phi(x_i, \theta_1^0)}{\pi_1^0 \phi(x_i, \theta_1^0) + \pi_2^0 \phi(x_i, \theta_2^0)}$$

Responsibility Interpretation

Formal Definition:

γ_{ij} is the **responsibility** of component j for data point i

Meaning:

- It is the **likelihood** that x_i is (purely) modeled by component ϕ_j
 - γ_{ij} represents the **portion of x_i** that is modeled (generated) by component ϕ_j
-

Hard vs. Soft Assignment

Hard Assignment

We can use γ_{ij} to determine $\delta_{ij} \rightarrow x_i$ can be assigned either to ϕ_1 or ϕ_2 (**maximum vote** $\rightarrow$ binarization)

$\rightarrow$ **K-means** type assignment: each data point is assigned to **one and only one** component (cluster)

Soft Assignment

EM is “softer”:

A data point **contributes** (via γ_{ij}) to **multiple components** (density modes)

EM calculates a **soft assignment** with:

$$\sum_j \gamma_{ij} = 1 \quad \forall i$$

Advantage: More robust and better reflects uncertainty in assignment

M-step (Maximization)

M-step Principle

Note: For 2 components, $\gamma_{i1} + \gamma_{i2} = 1$

Given the **responsibility** of each data point (γ_{ij}), we can estimate the parameters of each component by **weighted maximum likelihood**:

Parameter Estimation

Means:

$$\hat{\mu}_1 = \sum_{i=1}^N \frac{\gamma_{i1}}{\sum_{k=1}^N \gamma_{k1}} x_i$$

$$\hat{\mu}_2 = \sum_{i=1}^N \frac{\gamma_{i2}}{\sum_{k=1}^N \gamma_{k2}} x_i$$

Covariances:

$$\hat{\Sigma}_1 = \sum_{i=1}^N \frac{\gamma_{i1}}{\sum_{k=1}^N \gamma_{k1}} (x_i - \hat{\mu}_1)(x_i - \hat{\mu}_1)^T$$

$$\hat{\Sigma}_2 = \sum_{i=1}^N \frac{\gamma_{i2}}{\sum_{k=1}^N \gamma_{k2}} (x_i - \hat{\mu}_2)(x_i - \hat{\mu}_2)^T$$

Mixture Weights:

$$\pi_j = \frac{1}{N} \sum_{i=1}^N \gamma_{ij}$$

This formula ensures that $\sum_j \pi_j = 1$

EM Algorithm Convergence

Iterative Principle

The alternating **Expectation-Maximization** cycle increases the data likelihood with respect to the mixture model.

Convergence Criterion:

The process is iterated until **convergence**, i.e., when the data likelihood no longer changes (almost):

$$|LL(\theta^{t+1}, X) - LL(\theta^t, X)| < \epsilon$$

where ϵ is a tolerance threshold (e.g., $\epsilon = 10^{-6}$)

Monotonicity Property

Guaranteed Property: At each iteration, $LL(\theta^{t+1}, X) \geq LL(\theta^t, X)$

The log-likelihood **never decreases** (hill-climbing property)

EM Algorithm Limitations

Sensitivity to Initialization

Hill-Climbing: The algorithm depends on **initial parameters**

Consequences:

- **Sensitive to local maxima:** may converge to a local rather than global optimum
- Different initializations may yield different results

Initialization Strategies

Recommended Methods:

- **K-means:** use K-means clustering to initialize centers μ_j
- **K-means++:** improved K-means with smart initialization
- **Multiple Runs:** run the algorithm several times with different initializations and keep the best

Slow Convergence

Potentially slow convergence, depending on:

- Separation between components
 - Distribution shapes
 - Number of components
-

Generalization to c Components

General Latent Variables

Generalization:

$$\delta_{i1}, \delta_{i2}, \dots, \delta_{ij}, \dots, \delta_{ic}$$

where $\delta_{ij} = \text{true}$ if data point x_i is generated by component ϕ_j ($\delta_{ij} = \text{false}$ otherwise)

Responsibility:

γ_{ij} : expectation of δ_{ij} over all components

Parameters to Estimate:

$$\theta = [\pi_j, \theta_j]_{j \in [[c]]} = [\pi_j, [\mu_j, \Sigma_j]]_j$$

Complete EM Algorithm for c Components

Step 1: Initialization

Parameter Initialization:

$$\theta^0 = \{\pi_j^0, \mu_j^0, \Sigma_j^0\}_{j \in [[c]]}$$

Common Strategies:

- **General:**
 - $\pi_j = 1/c$
 - μ_j chosen randomly
 - $\Sigma_j = I_D$
- **Alternative:** use K-means for initialization

Step 2: E-step

Responsibility Calculation:

For each data point $i \in [[N]]$ and each component $j \in [[c]]$:

$$\gamma_{ij} = \frac{\pi_j \phi(x_i, \theta_j)}{\sum_{k=1}^c \pi_k \phi(x_i, \theta_k)}$$

Step 3: M-step

Mixture Parameter Estimation:

Means:

$$\mu_j = \frac{\sum_{i=1}^N \gamma_{ij} x_i}{\sum_i \gamma_{ij}}$$

Covariances:

$$\Sigma_j = \frac{X_j \Gamma_j X_j^T}{\text{Tr}(\Gamma_j)}$$

where $\Gamma_j = \text{diag}(\gamma_{1j}, \dots, \gamma_{Nj})$ and X_j is centered on μ_j

Weights:

$$\pi_j = \frac{\sum_i \gamma_{ij}}{N}$$

Step 4: Iteration

Iterate steps 2 and 3 until convergence

Practical Modeling

Parameters Influencing Convergence

The **a priori parameterization** of the mixture affects convergence:

Critical Parameters:

1. c : number of components
2. Σ_j : form of covariance matrices (diagonal, full, parameterized)

Risks:

- **Too flexible** or **too rigid** models $\rightarrow$ poor convergence or lack of convergence
 - **Over-parameterization:** $D \times D \times c + 2 \times c$ variables
If N is small, D large, and c large $\rightarrow$ **problematic over-parameterization**
-

Covariance Matrix Form

Over-Parameterization Problem

Dimension of Σ :

$$\Sigma \in \mathbb{R}^{D \times D}$$

Critical Example:

Character recognition with $x_i \in \mathbb{R}^{256}$

$\rightarrow$ Need to estimate $256^2 \times c$ parameters for covariance (with about 7000 data points)!

Parameterization Strategies

Spherical Models:

- $\Sigma = \sigma I_D$
- **1 parameter** per component

Diagonal Models:

- $\Sigma = \text{diag}(\sigma_1, \dots, \sigma_D)$
- **D parameters** per component

Full Models:

- $\Sigma \in \mathbb{R}^{D \times D}$
- $O(D^2)$ **parameters** per component

Practical Advice: Start with simple models (spherical) then increase complexity if needed

Preprocessing: Dimensionality Reduction with PCA**High-Dimensional Problem**

For high-dimensional data (e.g., $D = 256$), complex models do not converge because the number of parameters p is too large.

Solution: Prior PCA**Recommended Strategy:**

1. Apply **PCA** to reduce dimension from D to K dimensions
2. Apply **EM** in the space of the first K principal components

Practical Example:

- 50 principal components reconstruct 90% of the signal
- EM in the space of 50, 10, or 2 principal components

Advantages:

- Drastic reduction in the number of parameters
 - Noise elimination
 - Faster and more stable convergence
-

Selecting the Number of Components

Parsimony Principle

Problem:

The larger c , the fewer points can be assigned to each component (on average)

Objective:

Seek **parsimonious models**, i.e., with a **small number of parameters** to estimate

Bayesian Information Criterion (BIC)

Likelihood-Complexity Trade-off:

- The larger c , the better the likelihood estimation LL
- But the more complex the model (many parameters)

BIC Formula:

$$BIC(\theta, X) = -2 \cdot LL(\theta, X) + |\theta| \log(N)$$

where $|\theta|$ is the **number of parameters** in the model

Selection Criterion:

$$\hat{\theta} = \arg \min_{\theta} BIC(\theta, X)$$

BIC Components:

- $-2 \cdot LL(\theta, X)$: penalizes models with poor likelihood
- $|\theta| \log(N)$: penalizes model complexity

Interpretation: BIC favors models that explain data well **without overfitting**

Model Selection Limitations

Practical Challenges:

- Need to **test all models** (different values of c)
- Depends on **convergence** of each model
- Often requires **fine manual tuning**

Synthesis of Key Concepts

Gaussian Mixture Models (GMM)

Generalization:

The Gaussian mixture model **generalizes** the assumptions underlying PCA and LDA

Explicit Modeling and Estimation:

- **Explicit modeling** of density
- **Estimation** by maximum likelihood

Unsupervised Classification:

If components are classes, data point x_i is associated with class j for which:

$$\mathbb{P}(C = j|x_i) \approx \pi_j \phi_j(x_i)$$

is maximized among all classes

EM Algorithm

Characteristics:

- **Iterative algorithm** to maximize (log-)likelihood
- Principle used in **many other scenarios**
- Based on defining **unobserved (latent) variables** δ_{ij}

Applications:

- **Probabilistic Latent Semantic Analysis (pLSA)**: EM where hidden variables are latent concepts
 - **Hidden Markov Models (HMM)**
 - **Matrix Factorization**
-

Relationship Between GMM and K-means

Conceptual Comparison

K-means:

- **Hard assignment**
- Each point belongs to **one cluster only**
- Minimizes Euclidean distance to centers

GMM with EM:

- **Soft assignment**
- Each point has a **probability of belonging** to each component
- Maximizes likelihood according to a probabilistic model

GMM as a Generalization of K-means

K-means is a special case of GMM when:

- Covariance matrices $\Sigma_j = \sigma^2 I$ (spherical and identical)
 - $\sigma \rightarrow 0$: responsibilities γ_{ij} become binary
 - Soft assignment $\rightarrow$ hard assignment
-

Typical Exam Questions

Conceptual Questions

Gaussian Mixture Model:

- What is a Gaussian mixture model (GMM)?
- Why can it be seen as conditional modeling?
- How can it be used to generate data?

Responsibility:

- What is responsibility?
- How to interpret it?

EM Algorithm:

- Why is EM for GMM a MLE?
- How to interpret the E-step?
- How to interpret the M-step?

Technical Questions

Relationship with K-means:

- Discuss the relationship between GMM and K-means
- What are hard and soft assignments?

Convergence and Optimization:

- How to choose the number of components?
- What are the initialization strategies?
- How to handle over-parameterization?

Mathematical Developments

Important Advice: It is **strongly recommended** to develop the algebra contained in this chapter

Key Points to Master:

- Derivation of the log-likelihood
- Calculation of responsibilities using Bayes' rule
- Derivation of M-step update equations
- Proof of LL monotonicity in the EM algorithm

Mathematical Concepts to Master

Probability and Statistics

Fundamentals:

- Multivariate normal distribution
- Maximum likelihood estimation (MLE)
- Log-likelihood
- Bayes' rule

Modeling:

- Latent random variables
- Conditional probabilities
- Conditional expectation

Optimization**Techniques:**

- Iterative optimization
- Hill-climbing algorithms
- Convergence criteria
- Local vs. global maxima

Selection Criteria:

- BIC (Bayesian Information Criterion)
- Bias-variance trade-off
- Model parsimony