

Table of contents

The Essentials	2
Introduction and Motivation	2
The Fundamental Problem	2
Recap: Models Used So Far	3
New Approach: Density Modeling	3
Density Estimation Methods	3
Overview of Methods	3
Mixture Models: Definition	4
General Formulation	4
Model Assumptions	4
Probabilistic Interpretation	4
Mixture Models as “Humps” in the Landscape	4
Maximum Likelihood Estimation (MLE)	5
Likelihood Recap	5
Log-Likelihood	5
Gaussian Mixture Model (GMM)	6
Choice of Gaussian Family	6
Advantages of GMM	6
Simple Case: 1 Component	6
Direct Estimation	6
The Chicken-and-Egg Dilemma	7
Problem for 2 Components	7
The Vicious Circle	7
Latent Variables and EM Strategy	8
Introduction of Latent Variables	8
Binary Assignment Variables	8
The Two Magical Steps of EM	8
E-step (Expectation): Probabilistic Assignment	8
M-step (Maximization): Parameter Update	9
Hard vs. Soft Assignment	10
Hard Assignment	10
Soft Assignment	11
EM Algorithm Convergence	11
Iterative Principle	11
Monotonicity Property	11
EM Algorithm Limitations	12
Sensitivity to Initialization	12
Initialization Strategies	12
Slow Convergence	12

Complete EM Algorithm for c Components	12
Step 1: Initialization	12
Step 2: E-step	13
Step 3: M-step	13
Step 4: Iteration	13
Practical Modeling	14
Covariance Matrix Form	14
Parameterization Strategies	14
Preprocessing: Dimensionality Reduction with PCA	15
High-Dimensional Problem	15
Solution: Prior PCA	15
Selecting the Number of Components	15
The Complexity Dilemma	15
Bayesian Information Criterion (BIC)	16
Relationship Between GMM and K-means	16
Conceptual Comparison	16
GMM as a Generalization of K-means	17
Pitfalls to Avoid and Best Practices	17
Pitfall 1: “Ghost Components”	17
Pitfall 2: Premature Convergence	17
Pitfall 3: Abusive Interpretation	17
Synthesis of Key Concepts	18
Gaussian Mixture Models (GMM)	18
EM Algorithm	18

The Essentials

Introduction and Motivation

The **Expectation-Maximization (EM) algorithm** is a method for **estimating latent factors** in an **unsupervised** context of conditional modeling.

The Fundamental Problem

Imagine you’re in a room with several musical instruments playing simultaneously. You record the overall sound (the mixture). How can you **recover the characteristics of each instrument** (pitch, volume) without knowing who plays what and when?

This is exactly the problem that the EM algorithm solves for mixture models: **separating a mixture into its underlying components** when the origin of each observation is unknown.

The core of the problem:

With mixture models, we have **two types of unknowns** that depend on each other:

- **The parameters** of each component (mean, variance)
- **The origin** of each point (which component generated it)

This is a **vicious circle**: to know one, you need to know the other.

Recap: Models Used So Far

Component Models (PCA):

- Criterion based on the **variance** of centered data
- Implicit distribution: **normal**

Discriminant Models (LDA):

- Variance of projected data for within-class and between-class models

Common point: The distribution is **fixed** (essentially normal) and we seek its parameters θ (e.g., $\theta = [\mu, \Sigma]$).

New Approach: Density Modeling

Alternative: Seek a model for the **data density**.

Let $f_\theta(x) : \Omega \rightarrow \mathbb{R}$ be the **data density**.

Density Estimation Methods

Overview of Methods

Non-parametric Methods:

- **k-Nearest Neighbors (kNN):** Estimation based on the k nearest neighbors
- **Parzen Windows:** Kernel-based estimation
- **RBF Networks:** Radial basis functions
- **Histograms:** Space discretization

Parametric Methods:

- **Mixture Models:** The focus of this course
-

Mixture Models: Definition

General Formulation

The density $f(x)$ is generated by c “basis” functions ϕ (components), from a family $\mathcal{F}$:

$$f(x) = \sum_{j=1}^c \pi_j \phi(x, \theta_j)$$

Model Components:

- $\pi_j \in \mathbb{R}$: **mixture parameters**
- $\phi(x, \theta_j) \in \mathcal{F}$: **basis functions** controlled by parameters θ_j

Model Assumptions

Three Fundamental Assumptions:

- The **number of components** $c \in \mathbb{N}^*$ is **known**
 - The **function family** $\mathcal{F}$ is **known**
 - The **labels** (class labels) are **unknown** (unsupervised context)
-

Probabilistic Interpretation

Mixture Models as “Humps” in the Landscape

Imagine your data forms a **mountain landscape**. Each mountain represents a natural group in your data. A Gaussian mixture model (GMM) represents this landscape as a **superposition of Gaussian hills**.

The density $f(x)$ represents a **random process** where x is drawn from a set of states ω_j with prior probability $\mathbb{P}(\omega_j)$.

$$f(x) = \sum_{j=1}^c \pi_j \phi(x, \theta_j) = \mathbb{P}(x|\theta) = \sum_{j=1}^c \mathbb{P}(\omega_j) \mathbb{P}(x|\omega_j, \theta_j)$$

By Identification:

- $\pi_j = \mathbb{P}(\omega_j)$ with the constraint $\sum_j \pi_j = 1$

- $\phi(x, \theta_j) = \mathbb{P}(x|\omega_j, \theta_j)$ is the **conditional probability** that x is generated by state ω_j

Why Gaussians?

- **Central Limit Theorem:** Many natural phenomena are Gaussian
 - **Universality:** A mixture of Gaussians can approximate any density
 - **Computability:** The formulas are mathematically tractable
-

Maximum Likelihood Estimation (MLE)

Likelihood Recap

Let $X = \{x_1, \dots, x_N\}$ be **unlabeled** samples generated by the mixture $f(x) = \mathbb{P}(x|\theta)$.

Parameters to Infer:

$$\theta = [\pi_j, \theta_j]_{j \in [[c]]}$$

Assuming the data are **independent and identically distributed (i.i.d.)**:

$$\mathbb{P}(X|\theta) \stackrel{\text{i.i.d.}}{=} \prod_{i=1}^N \mathbb{P}(x_i|\theta)$$

Likelihood Estimator:

$$\hat{\theta} = \arg \max_{\theta} \mathbb{P}(X|\theta)$$

Log-Likelihood

Logarithmic Transformation:

$$LL(\theta, X) = \sum_{i=1}^N \log \mathbb{P}(x_i|\theta) = \sum_{i=1}^N \log \left[\sum_{j=1}^c \pi_j \phi(x_i, \theta_j) \right]$$

Maximum Log-Likelihood Estimator (MLE):

$$\hat{\theta} = \arg \max_{\theta} LL(\theta, X)$$

Advantage: Transforms the product into a sum, numerically more stable.

Gaussian Mixture Model (GMM)

Choice of Gaussian Family

Since ϕ must be a **probability density function**, choosing the **normal density family** as a basis seems reasonable (an **agnostic** choice):

$$\phi \in \mathcal{F} = \{\mathcal{N}(\mu, \Sigma)\}$$

Gaussian Mixture Model:

$$f(x) = \sum_{j=1}^c \pi_j \mathcal{N}(x | \mu_j, \Sigma_j)$$

where $\phi_j(x) := \phi(x, \theta_j) = \mathcal{N}(x | \mu_j, \Sigma_j)$.

Advantages of GMM

Universality: Can approximate **any density** (universal approximation theorem).

Linearity: Allows a **linear system** of parameters when maximizing the log-likelihood.

Simple Case: 1 Component

Direct Estimation

Parameters: $\theta = [\mu, \Sigma]$

i Mathematical Details: Direct Optimization

Optimization Problem:

$$\hat{\theta} = \arg \max_{\theta} \sum_i \log e^{-(x_i - \mu)^T \Sigma^{-1} (x_i - \mu)} \Leftrightarrow \hat{\theta} = \arg \min_{\theta} \sum_i (x_i - \mu)^T \Sigma^{-1} (x_i - \mu)$$

Analytical Solutions:

$$\hat{\mu} = \frac{1}{N} \sum_i x_i$$
$$\hat{\Sigma} = \frac{1}{N} \sum_i (x_i - \hat{\mu})(x_i - \hat{\mu})^T$$

Note: Solutions identical to standard empirical estimation of mean and covariance.

The Chicken-and-Egg Dilemma

Problem for 2 Components

Parameters:

$$\theta = [\pi_1, \theta_1, \pi_2, \theta_2] = [\pi_1, [\mu_1, \Sigma_1], \pi_2, [\mu_2, \Sigma_2]]$$

with the constraint $\pi_1 + \pi_2 = 1$.

Log-Likelihood:

$$LL(\theta, X) = \sum_{i=1}^N \log[\pi_1 \phi(x_i, \theta_1) + \pi_2 \phi(x_i, \theta_2)]$$

Difficulty: Maximization is challenging due to the **sum inside the logarithm!**

The Vicious Circle

To estimate the parameters (μ_j, Σ_j) , you need to know the origin of each point. To know the origin of each point, you need to know the parameters.

This is the classic “chicken-and-egg” problem!

The EM Solution: Instead of seeking a perfect solution all at once, EM proceeds by **successive approximations** by alternating two questions.

Latent Variables and EM Strategy

Introduction of Latent Variables

Fundamental Problem:

What is missing here is the **assignment** of x_i to one of the 2 components ϕ_j .

Key Insight:

If we **knew** this assignment, we would treat the problem as **two instances of the 1-component case**.

Binary Assignment Variables

We introduce **latent (unobserved) variables**: the binary assignment $\delta_{ij} \in \{\text{true}, \text{false}\}$ of each x_i to component ϕ_j :

- $x_i \sim \phi_1 \Rightarrow \delta_{i1} = \text{true}, \delta_{i2} = \text{false}$
- $x_i \sim \phi_2 \Rightarrow \delta_{i1} = \text{false}, \delta_{i2} = \text{true}$

But we can only **estimate** $\delta_{ij} \rightarrow$ **EM strategy**.

The Two Magical Steps of EM

E-step (Expectation): Probabilistic Assignment

Parameters:

$$\theta = [\pi_1, \theta_1, \pi_2, \theta_2]$$

Initialization:

We assume an initial value is known $\theta^0 = [\pi_1^0, \theta_1^0, \pi_2^0, \theta_2^0]$.

Responsibility Calculation:

We can infer the **statistical contribution** γ_{ij} of each data point x_i to each component ϕ_j :

$$\gamma_{ij}(\theta^0) = \mathbb{P}[\delta_{ij} = \text{true} | \theta^0, X]$$

i Mathematical Details: Bayes' Formula

Bayes' Formula:

$$\gamma_{i1}(\theta^0) = \frac{\pi_1^0 \phi(x_i, \theta_1^0)}{\pi_1^0 \phi(x_i, \theta_1^0) + \pi_2^0 \phi(x_i, \theta_2^0)}$$

Formal Definition:

γ_{ij} is the **responsibility** of component j for data point i .

Meaning:

- It is the **likelihood** that x_i is (purely) modeled by component ϕ_j
- γ_{ij} represents the **portion of** x_i that is modeled (generated) by component ϕ_j

General Interpretation:

For c components:

$$\gamma_{ij} = \frac{\pi_j \phi(x_i, \theta_j)}{\sum_{k=1}^c \pi_k \phi(x_i, \theta_k)}$$

Question: Given the current parameters, what is the probability that each point comes from each component?

This is a **soft assignment**. Unlike K-means which says “this point is 100% in this cluster,” EM says “this point is 60% in cluster A, 30% in B, 10% in C.”

M-step (Maximization): Parameter Update

Note: For 2 components, $\gamma_{i1} + \gamma_{i2} = 1$.

Given the **responsibility** of each data point (γ_{ij}), we can estimate the parameters of each component by **weighted maximum likelihood**.

i Mathematical Details: Parameter Estimation

Means:

$$\hat{\mu}_1 = \sum_{i=1}^N \frac{\gamma_{i1}}{\sum_{k=1}^N \gamma_{k1}} x_i = \frac{\sum_{i=1}^N \gamma_{i1} x_i}{\sum_{i=1}^N \gamma_{i1}}$$
$$\hat{\mu}_2 = \sum_{i=1}^N \frac{\gamma_{i2}}{\sum_{k=1}^N \gamma_{k2}} x_i = \frac{\sum_{i=1}^N \gamma_{i2} x_i}{\sum_{i=1}^N \gamma_{i2}}$$

Covariances:

$$\hat{\Sigma}_1 = \sum_{i=1}^N \frac{\gamma_{i1}}{\sum_{k=1}^N \gamma_{k1}} (x_i - \hat{\mu}_1)(x_i - \hat{\mu}_1)^T$$

$$\hat{\Sigma}_2 = \sum_{i=1}^N \frac{\gamma_{i2}}{\sum_{k=1}^N \gamma_{k2}} (x_i - \hat{\mu}_2)(x_i - \hat{\mu}_2)^T$$

Mixture Weights:

$$\pi_j = \frac{1}{N} \sum_{i=1}^N \gamma_{ij}$$

This formula ensures that $\sum_j \pi_j = 1$.

General Form for c Components:

$$\mu_j = \frac{\sum_{i=1}^N \gamma_{ij} x_i}{\sum_i \gamma_{ij}}$$

$$\Sigma_j = \frac{X_j \Gamma_j X_j^T}{\text{Tr}(\Gamma_j)}$$

where $\Gamma_j = \text{diag}(\gamma_{1j}, \dots, \gamma_{Nj})$ and X_j is centered on μ_j .

Question: Given the responsibilities, what are the best parameter estimates?

Parameter Recalculation:

- μ_j = weighted mean of points (weighted by γ_{ij})
- Σ_j = weighted variance of points
- π_j = average proportion of the component

Hard vs. Soft Assignment

Hard Assignment

We can use γ_{ij} to determine $\delta_{ij} \rightarrow x_i$ can be assigned either to ϕ_1 or ϕ_2 (**maximum vote** $\rightarrow$ binarization).

$\rightarrow$ **K-means** type assignment: each data point is assigned to **one and only one** component (cluster).

Soft Assignment

EM is “softer”:

A data point **contributes** (via γ_{ij}) to **multiple components** (density modes).

EM calculates a **soft assignment** with:

$$\sum_j \gamma_{ij} = 1 \quad \forall i$$

Advantage: More robust and better reflects uncertainty in assignment.

EM Algorithm Convergence

Iterative Principle

The alternating **Expectation-Maximization** cycle increases the data likelihood with respect to the mixture model.

Convergence Criterion:

The process is iterated until **convergence**, i.e., when the data likelihood no longer changes (almost):

$$|LL(\theta^{t+1}, X) - LL(\theta^t, X)| < \epsilon$$

where ϵ is a tolerance threshold (e.g., $\epsilon = 10^{-6}$).

Monotonicity Property

Guaranteed Property: At each iteration, $LL(\theta^{t+1}, X) \geq LL(\theta^t, X)$.

The log-likelihood **never decreases** (hill-climbing property).

Why does it work? Each step improves the **likelihood** of the data (probability of observing this data with the current model). It’s a “**hill-climbing**” algorithm: at each iteration, we climb the likelihood hill.

EM Algorithm Limitations

Sensitivity to Initialization

Hill-Climbing: The algorithm depends on **initial parameters**.

Consequences:

- **Sensitive to local maxima:** may converge to a local rather than global optimum
- Different initializations may yield different results

Initialization is crucial: EM can converge to **local maxima** (small peaks) instead of the global maximum (highest peak).

Initialization Strategies

Recommended Methods:

- **K-means:** use K-means clustering to initialize centers μ_j
- **K-means++:** improved K-means with smart initialization
- **Multiple Runs:** run the algorithm several times with different initializations and keep the best

Slow Convergence

Potentially slow convergence, depending on:

- Separation between components
- Distribution shapes
- Number of components

Complete EM Algorithm for c Components

Step 1: Initialization

Parameter Initialization:

$$\theta^0 = \{\pi_j^0, \mu_j^0, \Sigma_j^0\}_{j \in [c]}$$

Common Strategies:

- **General:** $\pi_j = 1/c$, μ_j chosen randomly, $\Sigma_j = I_D$
- **Alternative:** use K-means for initialization

Step 2: E-step

Responsibility Calculation:

For each data point $i \in [[N]]$ and each component $j \in [[c]]$:

$$\gamma_{ij} = \frac{\pi_j \phi(x_i, \theta_j)}{\sum_{k=1}^c \pi_k \phi(x_i, \theta_k)}$$

Step 3: M-step

Mixture Parameter Estimation:

$$\mu_j = \frac{\sum_{i=1}^N \gamma_{ij} x_i}{\sum_i \gamma_{ij}}$$

$$\Sigma_j = \frac{X_j \Gamma_j X_j^T}{\text{Tr}(\Gamma_j)}$$

where $\Gamma_j = \text{diag}(\gamma_{1j}, \dots, \gamma_{Nj})$ and X_j is centered on μ_j .

$$\pi_j = \frac{\sum_i \gamma_{ij}}{N}$$

Step 4: Iteration

Iterate steps 2 and 3 until convergence.

Practical Modeling

Covariance Matrix Form

Over-Parameterization Problem:

$$\Sigma \in \mathbb{R}^{D \times D}$$

Critical Example:

Character recognition with $x_i \in \mathbb{R}^{256} \rightarrow$ Need to estimate $256^2 \times c$ parameters for covariance!

Parameterization Strategies

Spherical Models:

$$\Sigma = \sigma I_D$$

$\rightarrow$ 1 parameter per component

Diagonal Models:

$$\Sigma = \text{diag}(\sigma_1, \dots, \sigma_D)$$

$\rightarrow D$ parameters per component

Full Models:

$$\Sigma \in \mathbb{R}^{D \times D}$$

$\rightarrow O(D^2)$ parameters per component

Practical Advice: Start with simple models (spherical) then increase complexity if needed.

Beware of over-parameterization! With $D=256$ dimensions and $K=10$ components, a full covariance requires 330,000 parameters! Often, dimension reduction with PCA is performed first.

Preprocessing: Dimensionality Reduction with PCA

High-Dimensional Problem

For high-dimensional data (e.g., $D = 256$), complex models do not converge because the number of parameters p is too large.

Solution: Prior PCA

Recommended Strategy:

- Apply **PCA** to reduce dimension from D to K dimensions
- Apply **EM** in the space of the first K principal components

Practical Example:

- 50 principal components reconstruct 90% of the signal
- EM in the space of 50, 10, or 2 principal components

Advantages:

- Drastic reduction in the number of parameters
 - Noise elimination
 - Faster and more stable convergence
-

Selecting the Number of Components

The Complexity Dilemma

Problem:

The larger c , the fewer points can be assigned to each component (on average).

Objective:

Seek **parsimonious models**, i.e., with a **small number of parameters** to estimate.

- **K too small:** Underfitting $\rightarrow$ model too simple, does not capture the structure
- **K too large:** Overfitting $\rightarrow$ model too complex, captures noise

Bayesian Information Criterion (BIC)

Likelihood-Complexity Trade-off:

- The larger c , the better the likelihood estimation LL
- But the more complex the model (many parameters)

BIC Formula:

$$BIC(\theta, X) = -2 \cdot LL(\theta, X) + |\theta| \log(N)$$

where $|\theta|$ is the **number of parameters** in the model.

Selection Criterion:

$$\hat{\theta} = \arg \min_{\theta} BIC(\theta, X)$$

BIC Components:

- $-2 \cdot LL(\theta, X)$: penalizes models with poor likelihood
- $|\theta| \log(N)$: penalizes model complexity

Interpretation: BIC favors models that explain data well **without overfitting** (balance between fit quality and complexity).

Relationship Between GMM and K-means

Conceptual Comparison

K-means:

- **Hard assignment**
- Each point belongs to **one cluster only**
- Minimizes Euclidean distance to centers
- Geometric model based on distances

GMM with EM:

- **Soft assignment**
- Each point has a **probability of belonging** to each component
- Maximizes likelihood according to a probabilistic model
- Probabilistic model based on densities

GMM as a Generalization of K-means

K-means is a special case of GMM when:

- Covariance matrices $\Sigma_j = \sigma^2 I$ (spherical and identical)
- $\sigma \rightarrow 0$: responsibilities γ_{ij} become binary
- Soft assignment $\rightarrow$ hard assignment

Key Point: K-means if you want simple and fast clustering with spherical clusters. EM if you want a probabilistic model, flexible shapes, or to handle uncertainty.

Pitfalls to Avoid and Best Practices

Pitfall 1: “Ghost Components”

Problem: EM can create components with very few points ($\pi_j \approx 0$) that capture noise.

Solution: Check the weights π_j after convergence, merge or eliminate very small components.

Pitfall 2: Premature Convergence

Problem: The algorithm seems to converge but remains in a poor local maximum.

Solution: Multiple trials with different initializations, monitor likelihood.

Pitfall 3: Abusive Interpretation

Problem: Thinking that the found components necessarily correspond to “real classes.”

Solution: EM finds **statistical groups**, not necessarily semantic categories. Validate with domain knowledge.

Synthesis of Key Concepts

Gaussian Mixture Models (GMM)

Generalization:

The Gaussian mixture model **generalizes** the assumptions underlying PCA and LDA.

Explicit Modeling and Estimation:

- **Explicit modeling** of density
- **Estimation** by maximum likelihood

Unsupervised Classification:

If components are classes, data point x_i is associated with class j for which:

$$\mathbb{P}(C = j|x_i) \approx \pi_j \phi_j(x_i)$$

is maximized among all classes.

EM Algorithm

Characteristics:

- **Iterative algorithm** to maximize (log-)likelihood
- Principle used in **many other scenarios**
- Based on defining **unobserved (latent) variables** δ_{ij}

Applications:

- **Probabilistic Latent Semantic Analysis (pLSA)**: EM where hidden variables are latent concepts
- **Hidden Markov Models (HMM)**
- **Matrix Factorization**
- Image segmentation, audio source separation

Condensed Cheat Sheet

Must Remember:

- EM = algorithm for models with latent variables

- Two alternating steps: E (responsibilities) $\rightarrow$ M (parameters)
- GMM = mixture of Gaussians, models data density
- Crucial initialization $\rightarrow$ use K-means or K-means++
- Choose K with BIC or cross-validation
- Covariance form: start simple (spherical), complexify if needed
- Main advantage: probabilistic model with soft assignments
- Relationship: generalizes K-means (special case)
- Applications: segmentation, source separation, advanced clustering
- Verification: always visualize and validate results

! Mathematical Concepts to Master

Probability and Statistics:

- Multivariate normal distribution
- Maximum likelihood estimation (MLE)
- Log-likelihood
- Bayes' rule
- Latent random variables
- Conditional probabilities
- Conditional expectation

Optimization:

- Iterative optimization
- Hill-climbing algorithms
- Convergence criteria
- Local vs. global maxima
- BIC (Bayesian Information Criterion)
- Bias-variance trade-off
- Model parsimony

Points to Develop:

- Derivation of the log-likelihood
- Calculation of responsibilities using Bayes' rule
- Derivation of M-step update equations
- Proof of LL monotonicity in the EM algorithm