

Table of contents

Detailed Course	2
Introduction and Motivation	2
General Context	2
Real Data Problem	2
Alternative Terminology	3
Overview of Clustering Methods	3
Hierarchical Agglomerative Clustering	4
Algorithm Principle	4
Distances Between Sets (Clusters)	4
k-means Strategy	4
Fundamental Assumptions	4
Formal Model	5
k-means Loss Function	5
Recap: Coordinate Descent Algorithm	6
Application to k-means	7
Assignment Update	7
Complete k-means Algorithm	8
Detailed Description	8
Algorithm Steps	8
Convergence Criteria	9
Process Visualization	10
Properties of the k-means Algorithm	11
Theoretical Guarantees	11
Geometric Interpretation	11
Link with the EM Algorithm	11
Sensitivity to Initialization	12
Variants, Links, and Relationships	12
k-means++: Smart Initialization	12
Relationship Between k-means and Gaussian Models	14
Intrinsic Link with the Euclidean Metric	14
Constrained Clustering	15
Synthesis of Key Concepts	16
Essential Characteristics of k-means	16
Comparison: k-means vs GMM vs Hierarchical Clustering	16
Choosing the Number of Clusters K	17
Typical Exam Questions	18
Conceptual Questions	18
Technical Questions	18
Mathematical Developments	19

Mathematical Concepts to Master	19
Algebra and Geometry	19
Optimization	19
Probability and Statistics	20
Practical Implementation Tips	20
Data Preprocessing	20
Algorithmic Parameters	20

Detailed Course

Introduction and Motivation

General Context

The **k-means** algorithm is a method for **estimating discrete latent factors** in an **unsupervised** clustering context.

Course Objectives:

- Understand the **geometric and statistical** properties of data
- Analyze data and develop tools for this analysis
- Specifically address the **(unsupervised)** approach to data clustering
- Understand the **assumptions** made in designing these tools
- Work on the **theory in depth**

Important Note: Clustering is similar to classification (unsupervised), density estimation, and is dual to outlier detection

Real Data Problem

Fundamental Observation:

Data generally does **not** come from a simple Gaussian process (variations around a mean prototype)

Characteristics of Real Data:

- Data distribution typically shows **non-uniformity**
- Presence of **regions of higher density**

Definition of Clustering:

Clustering is an **unsupervised** method that aims to discover **coherent groups** of data, corresponding to **density peaks** in the data

Alternative Terminology

Vector Quantization (VQ):

Often used synonym (in signal processing) for clustering, meaning “multi-dimensional quantization”

Overview of Clustering Methods

Main Methods

There are many clustering methods, including:

Hierarchical Methods:

- Hierarchical Agglomerative Clustering (HAC)

Partitioning Methods:

- **k-means** (Lloyd’s algorithm)
- k-medoids

Spectral Methods:

- Spectral clustering

Other Approaches:

- Community detection
- High-dimensional clustering
- Self-Organizing Maps (SOM)
- Neural Gas

Historical Note: Research on clustering has been active for at least 50 years

Hierarchical Agglomerative Clustering

Algorithm Principle

Iterative Process:

Step 1: Initialization

- Each data point is a cluster

Step 2: Search

- Find the closest pair of clusters

Step 3: Merge

- Merge these two clusters

Step 4: Iteration

- Repeat from step 2 until completion

Result: Produces a **dendrogram** of the data

Distances Between Sets (Clusters)

Possible Metrics:

- **Distance between centers** (centroid linkage)
- **Maximum distance** (complete linkage): $\max_{x \in C_1, y \in C_2} d(x, y)$
- **Minimum distance** (single linkage): $\min_{x \in C_1, y \in C_2} d(x, y)$

Limitation: The number of clusters is **unknown a priori**

k-means Strategy

Fundamental Assumptions

The **k-means** clustering algorithm assumes:

- A **normed metric space** (Euclidean) on the data
- An **unsupervised** assignment of data to clusters
- A **hard-assignment** algorithm: assignment is **binary**, each data point is assigned to **one and only one** cluster

Formal Model

Data:

Let $X = \{x_i\}_{i \in [[N]]} \subset \Omega \subseteq \mathbb{R}^D$, and given $K \in \mathbb{N}^*$

Model Variables:

- **Binary assignment variables (latent):**

$$Z = \{z_{ik}\} \quad \text{with} \quad z_{ik} \in \{\text{false}, \text{true}\} \equiv \{0, 1\} \quad \forall i, k$$

- **Cluster representatives:**

$$M = \{\mu_k\}_{k \in [[K]]} \quad \text{with} \quad \mu_k \in \mathbb{R}^D \quad \forall k$$

Complete Parameters:

$$\theta = [M, Z]$$

k-means Loss Function

Optimization Problem

k-means seeks the following assignment:

$$\hat{\theta} = [\hat{M}, \hat{Z}] = \arg \min_{\theta = [M, Z]} L(\theta, X)$$

Loss Function

Definition:

$$L(\theta, X) = \sum_{i=1}^N \sum_{k=1}^K z_{ik} \|x_i - \mu_k\|_2^2$$

Interpretation:

- Sum of **squared distances** between each point and its cluster representative
- The smaller this sum, the better the partition

Problem Complexity

Major Difficulty: Since the assignment is binary, exact optimization is **NP-Hard**

Practical Solution: Seek an **approximation via coordinate descent**

Recap: Coordinate Descent Algorithm

General Principle

Coordinate descent is an **alternating minimization algorithm**:

Problem:

Let $f : \mathbb{R}^D \rightarrow \mathbb{R}$, we seek:

$$x^* = \arg \min_{x \in \mathbb{R}^D} f(x)$$

Resolution Strategy

Definition of Partial Functions:

For $d \in [[D]]$, define $f^d : \mathbb{R}^D \times \mathbb{R} \rightarrow \mathbb{R}$:

$$f^d(x, y) = f([x(1), \dots, x(d-1), y, x(d+1), \dots, x(D)]^T)$$

Alternating Optimization:

Seek alternately the minimum along each coordinate:

$$x^{(t+1)}(d) = \arg \min_y f^d(x^{(t)}, y)$$

Principle: Fix all variables except one, optimize this variable, then move to the next

Application to k-means

Alternating Optimization of Z and M

We alternate the optimization of Z and M

Definitions of Partial Losses:

$$L_M(Z) = L(\theta, X)|_{M=M}$$

$$L_Z(M) = L(\theta, X)|_{Z=Z}$$

Optimization with Respect to M (Representatives)

Derivative Calculation:

$$\frac{\partial L_Z}{\partial \mu_k} = -2 \sum_{i=1}^N z_{ik} (x_i - \mu_k)$$

Optimality Condition:

The optimal representative M for a given assignment Z is reached when:

$$\frac{\partial L_Z}{\partial \mu_k} = 0 \quad \Rightarrow \quad \mu_k = \frac{\sum_{i=1}^N z_{ik} x_i}{\sum_{i=1}^N z_{ik}}$$

Fundamental Conclusion: Cluster representatives are their **centers of mass** (barycenters)

Assignment Update

Optimization with Respect to Z

Given a set of cluster representatives M , to minimize $L_M(Z)$:

Recall of the Loss:

$$L_M(Z) = \left[\sum_{i=1}^N \sum_{k=1}^K z_{ik} \|x_i - \mu_k\|_2^2 \right]_{M=M}$$

Assignment Rule

This is a **sum of positive terms** $\rightarrow$ assign:

$$z_{ik} = 1 \quad \text{only when} \quad k = \arg \min_m \|x_i - \mu_m\|_2^2$$

(and $z_{ik} = 0$ otherwise)

Conclusion: $L_M(Z)$ is minimized when **each data point is assigned to its nearest cluster representative**

Complete k-means Algorithm

Detailed Description

Input Data:

- $X = \{x_i\}_{i \in [N]}$ with $x_i \in \mathbb{R}^D$
- $K \in \mathbb{N}^*$ (number of clusters)

Algorithm Steps

Step 1: Initialization

Initialize cluster representatives $M^{(0)}$

Step 2: Assignment (E-step like)

Given cluster representatives $M^{(t)}$, assignment $Z^{(t)}$ associates each data point with its nearest cluster representative:

$$z_{ik}^{(t)} = \begin{cases} 1 & \text{if } k = \arg \min_m \|x_i - \mu_m^{(t)}\|_2^2 \\ 0 & \text{otherwise} \end{cases}$$

Step 3: Update (M-step like)

Given assignment $Z^{(t)}$, new cluster representatives $M^{(t+1)}$ are the centers of mass of data points assigned to clusters:

$$\mu_k^{(t+1)} = \frac{\sum_{i=1}^N z_{ik}^{(t)} x_i}{\sum_{i=1}^N z_{ik}^{(t)}}$$

Step 4: Convergence


Repeat from step 2 until convergence

Convergence Criteria

Convergence is reached when:

- The centers of mass **do not move much**: $\|\mu_k^{(t+1)} - \mu_k^{(t)}\| < \epsilon$
- Or the assignment is **stable**: $Z^{(t+1)} = Z^{(t)}$

Process Visualization



Important Note: Step 2 is similar to an **Expectation** step, and step 3 is similar to a **Maximization** step, considering a hard assignment (reference to the EM algorithm)

Properties of the k-means Algorithm

Theoretical Guarantees

Loss Decrease:

Since it performs **alternating minimization of convex functions**, it guarantees to **decrease the loss at each iteration**:

$$L(M^{(t+1)}, Z^{(t+1)}) \leq L(M^{(t)}, Z^{(t)})$$

Finiteness:

There is a **combinatorial but finite** number of possible assignments $\rightarrow$ the number of iterations is **(large but) finite**

Geometric Interpretation

Voronoi Diagram:

Step 2 is equivalent to constructing the **discrete Voronoi diagram** of X with centers M as seeds

Definition: The Voronoi diagram partitions space into regions, where each region contains all points closer to a given center than to any other center

Link with the EM Algorithm

Conceptual Comparison:

- **Step 2 (Assignment):** similar to an **Expectation** step (but hard)
- **Step 3 (Update):** similar to a **Maximization** step

Key Difference: k-means uses **hard assignment** while EM uses **soft assignment**

Sensitivity to Initialization

Problem of Local Optima

Major Limitation: Since exact optimization (optimal assignment) is **NP-Hard**, the k-means algorithm reaches a **local optimum**

Consequences:

- The quality of the result **varies with initialization**
- Different initializations can yield very different results

Practical Strategies

Randomization:

Randomization can be useful if computation is fast → run multiple times with different initializations

Heuristics:

Otherwise, use **heuristics for better initialization** (like k-means++)

Variants, Links, and Relationships

k-means++: Smart Initialization

Fundamental Principle

The principle of **k-means++** is to **maximally disperse** the initial cluster representatives $M^{(0)}$ over the data

Demonstrated Advantages:

1. Produces **higher quality** clusters
2. **Accelerates convergence** of k-means

Detailed k-means++ Algorithm

Step 1: First Center

Select the first cluster representative μ_1^0 **randomly** from the data X

Step 2: Distance Calculation

For all unselected data points x_i , calculate:

$$\Delta_i = \|x_i - \mu_m^0\|_2^2$$

the distance between x_i and its nearest representative μ_m^0 among the k representatives already selected

Step 3: Weighted Sampling

Sample the next representative μ_{k+1}^0 from X with **probability proportional** to the distribution Δ

$$\mathbb{P}(\text{choosing } x_i) = \frac{\Delta_i}{\sum_j \Delta_j}$$

Step 4: Iteration

Repeat from step 2 until K representatives are chosen

Algorithm Intuition

Key Idea: Points **far from already selected centers** have a higher chance of being chosen as new centers

Consequence: Initial centers are well **distributed in the data space**, covering different density regions

Practical Adoption

Widespread Use: This initialization strategy is used by several Data Science packages:

- MATLAB™
- Python SciKit Learn
- R
- And other standard libraries

Relationship Between k-means and Gaussian Models

k-means as a Special Case of GMM

Question: Why do we say that k-means considers a **Gaussian model** for clusters?

Answer:

k-means can be seen as a limiting case of GMM (Gaussian Mixture Model) where:

Model Assumptions:

- Each cluster follows a **normal distribution**
- **Spherical and identical** covariance matrices: $\Sigma_k = \sigma^2 I_D$ for all clusters
- Variance $\sigma^2 \rightarrow 0$: distributions become spikes

Consequence:

In this limit, the **soft assignment** of GMM becomes a **hard assignment**, and we recover k-means

Probabilistic Formulation

If we consider the probabilistic model:

$$\mathbb{P}(x|\mu_k, \sigma^2) = \frac{1}{(2\pi\sigma^2)^{D/2}} \exp\left(-\frac{\|x - \mu_k\|^2}{2\sigma^2}\right)$$

When $\sigma \rightarrow 0$, maximizing this probability is equivalent to minimizing $\|x - \mu_k\|^2$, which is exactly the k-means criterion

Intrinsic Link with the Euclidean Metric

Why k-means Depends on the Euclidean Metric?

Analysis of the Loss Function:

$$L(\theta, X) = \sum_{i=1}^N \sum_{k=1}^K z_{ik} \|x_i - \mu_k\|_2^2$$

The norm $\|\cdot\|_2$ is the **Euclidean norm**: $\|x\|_2^2 = \sum_d x(d)^2$

Practical Consequences

Scale Sensitivity:

k-means is **sensitive to scale**. Variables with large values dominate the distance

Solution: Normalize the data before applying k-means:

- Center: subtract the mean
- Scale: divide by the standard deviation

Limitation to Euclidean Spaces:

k-means cannot be directly applied to **non-Euclidean** data (graphs, texts, etc.)

Alternatives: For other metrics, use variants like k-medoids that work with arbitrary distances

Constrained Clustering

k-means Extension

Clustering can be **constrained** with:

MUST-LINK Constraints:

Two points must be in the **same cluster**

CANNOT-LINK Constraints:

Two points must be in **different clusters**

Practical Application

These constraints allow incorporating **prior knowledge** or **domain constraints** into the clustering process

Synthesis of Key Concepts

Essential Characteristics of k-means

Nature of the Method:

- k-means is part of the **unsupervised** family of data modeling techniques
- One of the **most popular** clustering techniques

Assignment:

- k-means performs **hard assignment**
- Each point belongs to **one cluster only**

Optimization:

- **Solid** but **NP-Hard** optimization criterion (loss)
- The k-means algorithm seeks an **approximate solution** via coordinate descent (alternating minimization)

Convergence:

- The loss **is not convex** → technique **sensitive to initialization**
- k-means++ is a **prior heuristic** for initialization, effective in practice

Extensions:

- Clustering can be **constrained** (MUST-LINK and CANNOT-LINK constraints)

Comparison: k-means vs GMM vs Hierarchical Clustering

Criterion	k-means	GMM (with EM)	Hierarchical Clustering
Assignment Type	Hard	Soft	Hierarchical
Cluster Shape	Spherical(same variance)	Elliptical(different covariances)	Flexible
Number of Clusters	Must be specified	Must be specified	Determined a posteriori(dendrogram)

Criterion	k-means	GMM (with EM)	Hierarchical Clustering
Complexity	$O(N \cdot K \cdot D \cdot T)$ where T is the number of iterations	$O(N \cdot K \cdot D^2 \cdot T)$	$O(N^2 \log N)$ or $O(N^3)$
Advantages	Fast simple	Probabilistic model varied cluster shapes	No need to specify K visualization via dendrogram
Disadvantages	Sensitive to initialization assumes spherical clusters	Slower more parameters to estimate	Expensive for large data no reassignment

Choosing the Number of Clusters K

Problem

Difficulty: The number of clusters K must be **specified a priori** in k-means

Selection Methods

Elbow Method:

Plot the loss L as a function of K and look for an “elbow” in the curve

Silhouette Score:

Measures clustering quality by comparing intra-cluster distance to inter-cluster distance

Gap Statistic:

Compares observed loss to expected loss under a reference distribution

Information Criteria:

- BIC (Bayesian Information Criterion)
- AIC (Akaike Information Criterion)

Practical Advice: Test several values of K and use multiple criteria

Typical Exam Questions

Conceptual Questions

Formal Description:

- Formally describe clustering
- In what sense is it an unsupervised technique?

Relationship with Metric:

- Explain why k-means is intrinsically linked to the Euclidean metric
- Why do we say that k-means considers a Gaussian model for clusters?

Accuracy:

- Is the k-means algorithm exact?
- What are the theoretical guarantees?

Technical Questions

Initialization:

- What are the principles for initializing k-means?
- Describe the k-means++ algorithm

Optimization:

- What is the coordinate descent algorithm?
- How does it apply to k-means?

Relationship with Other Methods:

- What is the relationship between k-means and GMM?
- What is the difference between hard and soft assignment?

Mathematical Developments

Important Advice: It is **strongly recommended** to develop the algebra contained in this chapter

Key Points to Master:

- Derivation of the loss function
 - Calculation of $\frac{\partial L_Z}{\partial \mu_k}$ and optimality condition
 - Proof that centers are barycenters
 - Proof of monotonic loss decrease
 - Relationship between k-means and limiting Gaussian model
-

Mathematical Concepts to Master

Algebra and Geometry

Fundamentals:

- Euclidean norm and distance
- Center of mass (barycenter)
- Voronoi diagram
- Metric spaces

Optimization

Techniques:

- Coordinate descent
- Alternating minimization
- Combinatorial optimization
- NP-Hard complexity

Properties:

- Convergence of iterative algorithms
- Local vs global optima
- Convex functions

Probability and Statistics

Concepts:

- Multivariate Gaussian distribution
 - Maximum likelihood
 - Mixture models
-

Practical Implementation Tips

Data Preprocessing

Recommended Steps:

1. **Normalization:** center and scale variables
2. **Outlier Management:** detect and handle outliers
3. **Dimensionality Reduction:** PCA if D is very large

Algorithmic Parameters

Initialization:

- Use **k-means++** by default
- Or run **multiple times** with random initializations

Convergence:

- Tolerance threshold: typically $\epsilon = 10^{-4}$
- Maximum number of iterations: avoid infinite loops

Choice of K:

- Test several values
- Use validation criteria