

UNIVERSITÉ DE GENÈVE

ANALYSE ET TRAITEMENT DE L'INFORMATION

14X026

TP A : Linear Algebra

Author: Donnet Michel Jean Joseph

E-mail: Donnet.Michel Jean Joseph@etu.unige.ch

September 2025



**UNIVERSITÉ
DE GENÈVE**

FACULTÉ DES SCIENCES
Département d'informatique

1 Theory questions

Use the slides and your notes to answer these questions. You do not need to write code. Provide short explanations and examples where possible.

1. What is a vector? What is a vector space? Give an example of a vector in $\mathbb{R}^2$ or $\mathbb{R}^3$.

A vector is a mathematical object that has a length, a direction and a sense.

A vector space is a space of vectors with linear operations such as addition between vectors and multiplication by a scalar. It can be generated by the combination of a few vectors, according to the dimension of the space. For instance, we can generate all vectors of a vector space of dimension 2 by using 2 vectors.

2. What do we mean by the "dimension" D of a dataset?

The dimension of a dataset indicates the minimum number of vectors useful to generate the entire dataset by combinations between these vectors.

3. Create two small datasets: one scalar dataset of 5 numbers, and one vector dataset of 5 vectors in $\mathbb{R}^3$. Compute their mean and variance. Then, write the formal mathematical definitions of mean and variance.

The dataset are the following:

$$A = \{1, 2, 3, 4, 5\} \in \mathbb{R}$$

$$B = \left\{ \begin{pmatrix} 3 & 9 & 6 \\ 8 & 4 & 5 \\ 5 & 6 & 6 \\ 0 & 8 & 5 \\ 6 & 3 & 2 \end{pmatrix} \right\} \in \mathbb{R}^3$$

- The mean is given by the formula 1

$$\mu = \frac{1}{n} \sum_{i=1}^n x_i \text{ for } \mathbf{x}_i, \mu \in \mathbb{R}^d \quad (1)$$

So for the first dataset, we have the following computation:

$$\mu = \frac{1}{n} \sum_{i=1}^n A_i = \frac{1+2+3+4+5}{5} = 3 \in \mathbb{R}$$

And for the second dataset, we have the following computation:

$$\mu = \frac{1}{n} \sum_{i=1}^n B_i = \frac{1}{n} (\sum_{i=1}^n B_{i1} \quad \sum_{i=1}^n B_{i2} \quad \sum_{i=1}^n B_{i3}) = \frac{1}{5} (22 \quad 30 \quad 24) = (4, 6, 4, 8) \in \mathbb{R}^3$$

- The variance σ is given by the formula 2 by taking the square racin of the result

$$\sigma_j^2 = \frac{1}{n} \sum_{i=1}^n (x_{ij} - \mu_j)^2 \text{ for } \mathbf{x}_i, \sigma \in \mathbb{R}^d \quad (2)$$

So we have for the first dataset

$$\sigma^2 = \frac{1}{n} \sum_{i=1}^n (A_i - \mu)^2 = \frac{1}{5} ((1-3)^2 + (2-3)^2 + (3-3)^2 + (4-3)^2 + (5-3)^2) = \frac{10}{5} = 2$$

So $\sigma = \sqrt{2} \in \mathbb{R}$.

And for the second dataset

– j = 1

$$\sigma_1^2 = \frac{1}{n} \sum_{i=1}^n (x_{i1} - \mu_1)^2 = \frac{1}{5} ((3-4.4)^2 + (8-4.4)^2 + (5-4.4)^2 + (0-4.4)^2 + (6-4.4)^2) = \frac{37.2}{5} = 7.44$$

– $j = 2$

$$\sigma_2^2 = \frac{1}{n} \sum_{i=1}^n (x_{i2} - \mu_2)^2 = \frac{1}{5} ((9-6)^2 + (4-6)^2 + (6-6)^2 + (8-6)^2 + (3-6)^2) = \frac{1}{5} (9+4+0+4+9) = \frac{26}{5} = 5.2$$

– $j = 3$

$$\sigma_3^2 = \frac{1}{n} \sum_{i=1}^n (x_{i3} - \mu_3)^2 = \frac{1}{5} ((6-4.8)^2 + (5-4.8)^2 + (6-4.8)^2 + (5-4.8)^2 + (2-4.8)^2) = \frac{10.8}{5} = 2.16$$

So

$$\sigma = (\sqrt{7.44} \quad \sqrt{5.2} \quad \sqrt{2.16}) \in \mathbb{R}^3$$

4. Define the scalar product. Compute it for two vectors from your vector dataset above. How does the scalar product allow us to compute the projection of one vector onto another?

The scalar product is defined for two vectors $\mathbf{x}$ and $\mathbf{y}$ like the formula 3

$$\mathbf{x} \cdot \mathbf{y} = \sum_{i=1}^d x_i y_i = |\mathbf{x}| |\mathbf{y}| \cos(\theta) \quad \forall \mathbf{x}, \mathbf{y} \in S^d \quad (3)$$

with S the vector space of dimension d .

So for vector $\mathbf{x} = (3 \quad 9 \quad 6)$ and $\mathbf{y} = (8 \quad 4 \quad 5)$, the scalar product is given by the following calculs.

$$\mathbf{x} \cdot \mathbf{y} = \sum_{i=1}^d x_i y_i = 3 \times 8 + 9 \times 4 + 6 \times 5 = 90$$

5. What is a basis of a vector space?

A basis of a vector space is a set of vectors that can generate the entire vector space by basics operations, such as addition between vectors and multiplication by scalars.

6. Explain the eigenvalue decomposition (EVD) and singular value decomposition (SVD). Why are they important in data science?

$X \in \mathbb{R}^{n \times n}$ is a square matrix. The eigen vector $\mathbf{u}$ is a vector that satisfies the equality 4.

$$(\exists \sigma \text{ such as } X\mathbf{u} = \sigma\mathbf{u} \text{ with } \mathbf{u} \neq 0) \Rightarrow (X - \sigma I)\mathbf{u} = 0 \quad (4)$$

So if X is diagonalisable, the matrix admits an Eigen Value Decomposition, given by the formula 5.

$$\forall i, X\mathbf{u}_i = \sigma_i \mathbf{u}_i \text{ and } \mathbf{u}_i \neq 0 \Rightarrow XU = U\Sigma \Leftrightarrow X = U\Sigma U^{-1} \quad (5)$$

Where

- $U \in \mathbb{R}^{n \times n}$ is a matrix composed of linearly independant eigen vectors $\mathbf{u}$
- $\Sigma \in \mathbb{R}^{n \times n}$ is a diagonal matrix with on the diagonal the values of the corresponding eigenvalues σ_i of the eigenvectors.
- U^{-1} the inverse of U .

Similarly, $\forall X \in \mathbb{R}^{n \times m}$, there exists a factorization $X = USV^T$ which is called the Singular Value Decomposition (SVD).

This two notions are very important in data science because they allow us to uncorrelate data, which is very useful for data analysis and processing because it saves a lot of time and ressources.

7. Give one concrete example of the curse of dimensionality, and one example of the blessing of dimensionality.
-

- Curse of dimensionality. A concrete example of this can be the data sparsity problem. We have an hypercube $[a, b]^D$ with D the dimensionality of the hypercube. We want to create an histogram that represents an empiric estimation of the probability density function (pdf) for this hypercube. For that, we quantize each dimension into b bins. Now, we hope an average of n samples per bin since if the number of samples per bin is too small, we can't have a good estimation because the density can vary a lot between bins. This generate noise for our estimation whereas with a great number of samples per bin, the variance will decrease and the empiric estimation will be closest to the real pdf, according to the Central Limit Theorem which says that when the number of samples increase, the empiric mean is closest to the real mean.
- Blessing of dimensionality. According to the lemma of Johnson-Lindenstrauss, we can find a linear projection that maps the datas in a space of smaller dimensionality that conserves the structure of our data with a distortion factor $\epsilon \in]0, 1[$. This lemma is a blessing of dimensionality because it says also that the projection depends only on the number of samples and not on the dimensionality of our data. So for instance, we have 10000 samples $\in \mathbb{R}^{3000}$ and we choose a distortion factor ϵ of 50% According to the lemma, we can project this datas with linear projection into a space of dimensionality given by the formula 6.

$$d = O\left(\frac{\log_2 N}{\epsilon^2}\right) = O\left(\frac{\log_2 10000}{0.5^2}\right) = O\left(\frac{13.28}{0.25}\right) \approx 50 \quad (6)$$

So we can reduce the dimensionality from 3000 to only 50 !!! This is an answer to the curse of dimensionality since it allows to work with big dimensionalities.

8. If you divide a D -dimensional space into b bins per dimension, and want k samples per bin, how many samples are needed in total?

If we divide each dimension into b bins, we have for D dimensions a number of bins:

$$\underbrace{b \times b \times \dots \times b}_{D \text{ times}} = b^D$$

Now, if we want k samples per bin, the total number of samples needed is:

$$N = k \times b^D$$

9. Why does data become more sparse as D grows? What happens if the number of samples stays fixed?

We suppose that we have a fixed number of samples N . The number of samples per bin is given by:

$$N = k \times b^D \Leftrightarrow k = \frac{N}{b^D}$$

So the expectation to have a sample x_i in a bin b_j is given by:

$$\mathbb{E}[P(x_i \in b_j)] = \frac{1}{N} \cdot \frac{N}{b^D} = \frac{1}{b^D}$$

When D increase, this expectation will approach 0, so it means that when D increase, most bins are likely to be empty since expectation to have a sample inside a bin approach 0.

10. Compare a hypersphere of radius r with its enclosing hypercube $[-r, r]^D$. As D increases, what happens to their volume ratio? What does this imply for uniform sampling?

We have the volum of the hypersphere of dimension D and radius r given by the formula 7.

$$\text{vol}(B^D(r)) = \frac{2r^D \pi^{D/2}}{D \Gamma(D/2)} \quad (7)$$

And the volum of the hypercube of $[-r, r]^D$ given by the formula 8.

$$\text{vol}(C^D(r)) = (2r)^D \quad (8)$$

So the ratio of this two volumes is given by the formula 9.

$$\frac{\text{vol}(B^D(r))}{\text{vol}(C^D(r))} = \frac{2r^D \pi^{D/2}}{D \Gamma(D/2)} \cdot \frac{1}{(2r)^D} = \frac{\pi^{D/2}}{D \Gamma(D/2) \cdot 2^{D-1}} \quad (9)$$

When D increase, we have that:

$$\lim_{D \rightarrow \infty} \frac{\text{vol}(B^D(r))}{\text{vol}(C^D(r))} = \lim_{D \rightarrow \infty} \frac{\pi^{D/2}}{D \Gamma(D/2) \cdot 2^{D-1}} = 0$$

So when the dimension is big, the volum of the hypersphere is neglectible compared to the volum of the hypercube. So this implies for uniform sampling that if we make uniform sampling in hypercube, there are no samples in hypersphere because the ratio of volumes is 0. It means that all samples will be concentrated, like the volum, in the corners of the hypercube and not in the hypersphere.

11. In high dimensions, what is the typical angle between two random vectors? How does this angle change as D increases?

We create two vectors by creating 4 random variables independant and identically distribued (iid) W, X, Y, Z in the space $\Omega \subset \mathbb{R}^D$ of dimension D . We define $u = X - Y$, and $v = Z - W$ two vectors. Now, we are interested to compute the $\mathbb{E}(\cos(\theta))$ between the two vectors. We have:

$$\mathbb{E}[\cos(\theta)] = \mathbb{E}\left[\frac{\langle u, v \rangle}{\|u\| \cdot \|v\|}\right] = \mathbb{E}\left[\frac{\langle X - Y, Z - W \rangle}{\|X - Y\| \cdot \|Z - W\|}\right]$$

As W, X, Y, Z are iid, $X - Y$ and $Z - W$ are also iid and centered. So $\mathbb{E}[X - Y] = \mathbb{E}[Z - W] = 0$ so we have $\mathbb{E}[\langle X - Y, Z - W \rangle] = 0$

So $\mathbb{E}[\theta] = 0 \Rightarrow \theta = \frac{\pi}{2}$

12. What is meant by a “concentration phenomenon”? Give one example where it is useful, and one where it is problematic.

Write your answer here

13. State one concentration inequality (e.g., Markov, Chebyshev, Hoeffding). Explain briefly what it tells us.

Write your answer here

14. What is a PAC (Probably Approximately Correct) analysis in machine learning?

Write your answer here

2 Matrix

1. Find a quadratic polynomial, say $f(x) = ax^2 + bx + c$, such that

$$f(1) = -5, \quad f(2) = 1, \quad f(3) = 11.$$

Write your solution here

2. Let a, b be some fixed parameters. Solve the system of linear equations

$$\begin{cases} x + ay = 2 \\ bx + 2y = 3 \end{cases}$$

Write your solution here

3. Find the inverse of the following matrix, if it exists:

$$A = \begin{pmatrix} 1 & 2 & 3 \\ 0 & 1 & 4 \\ 5 & 6 & 0 \end{pmatrix}$$

Write your solution here

4. Solve the following matrix equation for (X):

$$AX = B, \quad \text{where} \quad A = \begin{pmatrix} 1 & 3 \\ 2 & 5 \end{pmatrix}, \quad B = \begin{pmatrix} 7 & 8 \\ 9 & 10 \end{pmatrix}.$$

Write your solution here

3 The importance of the mathematical concept behind a code

1. Download, open and run several times the PYTHON file “some_script.py”.

Explain the function `def project_on_first(u, v)` in geometrical terms and then explain the results of this code with clear mathematical concepts (you have to talk about projections and scalar product).

Write your solution here

2. Consider $u = (u_i)_{1 \leq i \leq 7}$ and $v = (v_i)_{1 \leq i \leq 7}$, two vectors of length 7. How would you clearly write the following part of the script using some mathematical notation?

```

1  import numpy as np
2  u = np.random.randn(7)
3  v = np.random.randn(7)
4  r = 0
5  for ui, vi in zip(u, v):
6      r += ui * vi

```

Name the mathematics behind and rewrite the last 3 lines of code in one.

Write your solution here

3. Complete the code to construct a vector v orthogonal to the vector u and with the same norm as u . Comment each line of your code.

Write your solution here

4. Given two vectors (u) and (v) in $\mathbb{R}^n$, write Python code to compute the cosine of the angle between them. Explain the mathematical concept behind this computation, and show how you would write it in one equation.

Write your solution here

4 Computing Eigenvalues, Eigenvectors, and Determinants

1. Find the determinant, eigenvectors, and eigenvalues of the matrix

$$\begin{bmatrix} 5 & 6 & 3 \\ -1 & 0 & 1 \\ 1 & 2 & -1 \end{bmatrix}$$

Compare your results with those of the computer.

Write your solution here

2. The covariance matrix for n samples $x_1, \dots, x_n$, each represented by a $d \times 1$ column vector, is given by

$$C = \frac{1}{n-1} \sum_{i=1}^n (x_i - \mu)(x_i - \mu)^T,$$

where C is a $d \times d$ matrix and $\mu = \sum_{i=1}^n x_i / n$ is the sample mean.

Prove that C is always positive semidefinite. (Note: A symmetric matrix C of size $d \times d$ is positive semidefinite if $v^T C v \geq 0$ for every $d \times 1$ vector v .)

Write your solution here

3. In this portion of the exercise, we will calculate the eigenvalues of the covariance matrices of six data sets listed as follows:

filename	n	d	description
tp1_artificialdata[1-3]	1024	100	Artificial data generated from various auto-regression
tp1_artificialdata4	1024	100	Random Gaussian data
tp1_freyfaces	1965	560	Facial images of a man named Brendan
tp1_digit2	5958	784	Hand-written images of "2"

Download and unzip data for python.zip. Each file contains a $n \times d$ data matrix with rows representing n different samples in $\mathbb{R}^d$.

- Compute the covariance matrix of each dataset.
- Compute the eigenvalues for each covariance matrix.
- Compute the determinant of each covariance matrix.
- Compute the product of the eigenvalues for each covariance matrix.
- Display the eigenspectrum for each covariance matrix in a 2D plot.

Describe the relationship between the product of the eigenvalues and the determinant of each covariance matrix. In addition, describe the observations you made regarding the plot of the spectrum of eigenvalues of each covariance matrix. How can you explain the disparities between datasets?

Write your solution here

4. Let $A \in \mathbb{R}^{n \times n}$ be a square matrix. Show that if (A) has an eigenvalue $(\lambda = 0)$, then the matrix (A) is singular. Furthermore, prove that the determinant of (A) must be zero if any eigenvalue is zero.

Write your solution here

5 Computing Projection Onto a Line

There are given a line $\alpha : 3x + 4y = -6$ and a point A with the coordinates $(-1, 3)$.

1. Find a distance from the point A to the line α using linear algebra (not coordinate method).
2. Explain your code using a sketch and some mathematics (there should be a scalar product somewhere).

`_Write your solution here_`